



SAPIENZA
UNIVERSITÀ DI ROMA

“SAPIENZA” UNIVERSITY OF ROME
FACULTY OF INFORMATION ENGINEERING,
INFORMATICS AND STATISTICS
DEPARTMENT OF COMPUTER SCIENCE

Machine Learning

Lecture notes integrated with the book "Machine Learning", Tom Mitchell

Author
Simone Bianco

January 27, 2025

Contents

Information and Contacts	1
1 Introduction on machine learning	2
1.1 What is machine learning?	2
1.2 Hypotheses, consistency and representation	6
1.2.1 Representation power and generalization power	7
1.3 Performance evaluation	9
1.3.1 Hypothesis comparison	11
1.3.2 Performance metrics	12
2 Classification problems	16
2.1 Decision trees learning	16
2.1.1 Entropy and the ID3 algorithm	18
2.1.2 Overfitting in decision trees	21
2.2 Bayesian learning	24
2.2.1 Uncertainty and probability	24
2.2.2 Optimal Bayes classifier	29
2.2.3 Naïve Bayes classifier	34
2.3 Probabilistic models for classification	37
2.3.1 Probabilistic generative models	37
2.3.2 Probabilistic discriminative models	39
2.4 Linear models for classification	41
2.4.1 Least squares method	42
2.4.2 Perceptron	43
2.4.3 Fisher's linear discriminant	45

Information and Contacts

Personal notes and summaries collected as part of the *Machine Learning* course offered by the degree in Computer Science of the University of Rome "La Sapienza".

Further information and notes can be found at the following link:

<https://github.com/Exyss/university-notes>. Anyone can feel free to report inaccuracies, improvements or requests through the Issue system provided by GitHub itself or by contacting the author privately:

- Email: bianco.simone@outlook.it
- LinkedIn: [Simone Bianco](#)

The notes are constantly being updated, so please check if the changes have already been made in the most recent version.

Suggested prerequisites:

Sufficient knowledge of calculus, probability and algorithm design

Licence:

These documents are distributed under the [GNU Free Documentation License](#), a form of copyleft intended for use on a manual, textbook or other documents. Material licensed under the current version of the license can be used for any purpose, as long as the use meets certain conditions:

- All previous authors of the work must be **attributed**.
- All changes to the work must be **logged**.
- All derivative works must be **licensed under the same license**.
- The full text of the license, unmodified invariant sections as defined by the author if any, and any other added warranty disclaimers (such as a general disclaimer alerting readers that the document may not be accurate for example) and copyright notices from previous versions must be maintained.
- Technical measures such as DRM may not be used to control or obstruct distribution or editing of the document.

1

Introduction on machine learning

1.1 What is machine learning?

While many common tasks can be easily solved by computers through an algorithm, some are hard to formalize as a series of steps to be executed in a deterministic way. As an analogy, consider how language is made of syntax and semantics. Syntax can easily be formalized as a sequence of sub-structures that make up a phrase. If a sentence is slightly malformed, the machine can have an hard time trying to reconstruct the correct syntax, but in most cases this can be achieved. For semantics, instead, we have a whole different problem: some words could have more meanings, giving sentences different interpretations depending on the context of the conversation. This task is clearly harder for a machine. Sometimes, not even humans are capable of solving it!

In the past, these type of tasks were solved through *expert systems*, that being any system programmed by an human expert to solve a specific task. Expert systems can be viewed as a sequence of if-else conditions: if the task requires x then do y , and so on. Not all tasks can be solved through expert systems. In particular, some tasks need different solutions for many cases, making these primitive systems useless due to all cases being impossible to program.

To solve this type of complex and variable tasks, we use **machine learning**, which slowly teaches the machine how to solve the problem in the best way possible. The idea here is to program computers in a way that improves a specific *performance criterion* through *example data* and *past experiences*.

Machine learning uses **data mining** — the act of producing knowledge from known data — to increase the experience of the machine in solving the designed problem. In general, machine learning comes in handy when one of the following conditions holds:

- There is no human expertise on the task
- Human experts are unable to explain their expertise
- The solution needs to adapt to particular cases

The field of machine learning had an exponential growth in recent years due to the growing flood of online data — the so called *big data phenomenon* — and the increase of computational power to process such data through advanced algorithms based on theoretical results. First, we give a formal definition of a *learning problem*.

Definition 1: Learning problem

A **learning problem** is the improvement over a task T with respect to a performance measure P based on experience E .

For example, suppose that we want to program a machine that learns how to play checkers. We define the learning problem as:

- The task T is to play checkers
- The performance measure P is the percentage of games won in a tournament
- The experience E is the opportunity to play against self

But how can we improve such performance measure? What *exactly* should the machine learn? These questions reduce the learning problem to finding a valid mathematical representation of T , P and E . The training process can be described by four phases:

1. The human expert suggests what is an optimal move for each configuration of the board
2. The human expert evaluates each configuration, ranking them by optimality
3. The computer plays against an human an automatically detects with configurations lead to a win, a loss or a draw
4. The computer plays against itself to improve performance

Formally, this whole process can be expressed as a simple mathematical function called **target function**. In particular, we want to choose a target function that represents the learning problem in the best way possible and that can be computed by a machine.

For instance, consider the function $V : \text{Board} \rightarrow \mathbb{R}$, defined as follows:

- If b is a final board state and it is a win then $V(b) = 100$
- If b is a final board state and it is a loss then $V(b) = -100$
- If b is a final board state and it is a draw then $V(b) = 0$
- If b is not a final board state then $V(b) = V(b^*)$, where b^* is the best final board state that can be achieved starting from the board b playing the optimal moves

This function perfectly models our learning problem. However, it cannot be computed by any program since we haven't defined what an optimal set of moves is. We need a new definition that encodes this concept of optimal strategy for a checkers game.

For example, we can re-define V as follows:

$$V(b) = w_0 + w_1 \cdot \text{bp}(b) + w_2 \cdot \text{rp}(b) + w_3 \cdot \text{bk}(b) + w_4 \cdot \text{rk}(b) + w_5 \cdot \text{bt}(b) + w_6 \cdot \text{rt}(b)$$

where:

- $\text{bp}(b)$ is the number of black pieces
- $\text{rp}(b)$ is the number of red pieces
- $\text{bk}(b)$ is the number of black kings
- $\text{rk}(b)$ is the number of red kings
- $\text{bt}(b)$ is the number of red pieces threatened by black pieces
- $\text{rt}(b)$ is the number of black pieces threatened by red pieces

With this formulation, we have reduced the concept of leaning checkers to estimating the best possible values of the coefficients w_1, \dots, w_6 , which are called **weights**, that maximize the value of $V(b)$ for any board state b . This estimation process is referred to as *learning the function V* .

Definition 2: Learned function

Given a target function $f : X \rightarrow Y$ with weights w_1, \dots, w_k , we define the **learned function** $\hat{f} : X \rightarrow Y$ as the current approximation of f computed by a learning algorithm.

By definition, the learned function \hat{f} will never be equal to the target function f : the target function's weights are always unknown by definition. The idea here is to approximate f by repeatedly applying small changes to the weights $\hat{w}_1, \dots, \hat{w}_k$ of \hat{f} in order to estimate the weights of f . To learn a function f , we need a *dataset*. A dataset is a set of instances that can be used by a learning algorithm to improve the performance of the learned function.

Definition 3: Dataset

Let $f : X \rightarrow Y$ be a target function and let $f_{\text{train}}(x)$ be the training value obtained by x in the training data. Given a set of n training inputs $X_D = \{x_1, \dots, x_n\}$, the **dataset** of the learning problem is the set of samples defined as:

$$D = \{(x_i, f_{\text{train}}(x_i)) \mid x_i \in X_D\}$$

After training, the learned function \hat{f} will have learned the values of the inputs in the dataset, returning a value as close as possible to the one in the dataset (in some cases the returned value is exactly the same). However, we are interested in estimating the *other* possible inputs, i.e. those that aren't in the dataset.

In summary, a machine learning problem is the task of learning a target function $f : X \rightarrow Y$ through a dataset D for a set X_D of n inputs. To learn a function f means computing an approximating function \hat{f} that returns values as close as possible to f , especially for values outside of the dataset D , implying that $\forall x \in X - X_D$ it should hold that $f(x) \approx \hat{f}(x)$. In order for the learned function to be *good*, the set of training inputs X_D must be very very small compared to the set of total inputs, meaning that $|X_D| \lll |X|$.

There are distinct types of machine learning problems based on:

- The type of dataset used:
 1. **Supervised learning**: problems where the model learns patterns from labeled data. Here, the dataset corresponds to $D = \{(x_i, y_i) \mid i \in X_D\}$, where y_i is the sample of the function value for x_i
 2. **Unsupervised learning**: problems where the model learns patterns from unlabeled data. Here, the dataset corresponds to $D = \{x_i \mid i \in X_D\}$
 3. **Reinforcement learning**: problems in which an agent learns to make decisions by interacting with an environment and receiving rewards or penalties based on its actions.
- The type of function to be learned:
 1. **Discrete Classification**: the input set is $X = A_1 \times \dots \times A_n$, where A_i is a finite set, and the output set is $Y = \{C_1, \dots, C_k\}$. When $k = 2$, i.e. we have only two classes, we say that the problem is a **Concept Learning** problem
 2. **Discrete Regression**: the input set is $X = A_1 \times \dots \times A_n$, where A_i is a finite set, and the output set is $Y = \mathbb{R}^m$.
 3. **Continuous Classification**: the input set is $X = \mathbb{R}^n$ and the output set is $Y = \{C_1, \dots, C_k\}$.
 4. **Continuous Regression**: the input set is $X = \mathbb{R}^n$ and the output set is $Y = \mathbb{R}^m$.

Classification problems are based on the classification of inputs into predetermined categories, while regression problems involve the approximation of functions defined over \mathbb{R} . Reinforcement learning, instead, is used for dynamic systems with unknown or partially known evolution model, usually robotic tasks and game playing.

1.2 Hypotheses, consistency and representation

After discussing the basic notation and terminology, we are ready to deepen our understanding on how to learn a problem.

Definition 4: Hypothesis space

Given a target function f , an **hypothesis space** H for f is a set of functions $h \in H$, where h is called hypothesis, that can be learned in order to reach an approximation of f .

The representation of an hypothesis space highly depends on the type of problem. For instance, consider the problem of classifying natural numbers into primes numbers and composite numbers. This corresponds to a discrete classification problem with two classes \mathbb{P} and $\mathbb{N} - \mathbb{P}$, i.e. a concept learning problem where we want to learn the concept of prime number. The target function of the problem is thus described as $f : \mathbb{N} \rightarrow \{\mathbb{P}, \mathbb{N} - \mathbb{P}\}$. Here, the simplest hypothesis space is the set $H = \{h_A : \mathbb{N} \rightarrow \{A, \mathbb{N} - A\} \mid A \subseteq \mathbb{N}\}$.

Given a performance measure P over a dataset D and an hypothesis space H , the learning task is to find the best approximation $h^* \in H$ of the function f using the dataset D .

$$h^* \in \arg \max_{h \in H} P(h, D)$$

Finding such optimal hypothesis is the core of a learning problem. However, by definition, an hypothesis space may also contain hypotheses that are clearly a wrong approximation of f over D . For instance, given a dataset D for a target function f , for any hypothesis $h \in H$ we can check whether $h(x) = f(x)$ only for instances $x \in X_D$ since the other values are unknown in D . This means that some hypothesis may have some values that are *inconsistent* with the dataset itself, making them useless in the learning process. An hypothesis $h \in H$ is said to be **consistent** with a dataset D of a target function f if and only if $h(x) = f(x)$ for all $x \in X_D$. The subset of hypothesis that are consistent with a dataset is called **version space**.

Definition 5: Version space

The **version space** of a target function f with respect to the hypothesis space H and the dataset D , written as $VS_{H,D}$, is the subset of H that contains all the hypotheses that are consistent with D .

$$VS_{H,D} = \{h \in H \mid \forall x \in X_D \ h(x) = f(x)\}$$

Consider again the previous example. Suppose that we're working with the following dataset $D = \{(1, \mathbb{N} - \mathbb{P}), (3, \mathbb{P}), (5, \mathbb{P}), (6, \mathbb{N} - \mathbb{P}), (7, \mathbb{P}), (10, \mathbb{N} - \mathbb{P})\}$. Here, the version space would restrict our interest to all those functions $h_A : \mathbb{N} \rightarrow \{A, \mathbb{N} - A\}$ where the elements 3, 5, 7 lie inside A and the elements 1, 6, 10 lie outside of A .

$$VS_{H,D} = \{h_A : \mathbb{N} \rightarrow \{A, \mathbb{N} - A\} \mid \{3, 5, 7\} \subseteq A \subseteq \mathbb{N} \text{ and } \{1, 6, 10\} \subseteq \mathbb{N} - A\}$$

By definition, the best approximation for a performance measure P over a dataset D and an hypothesis space H must clearly lie inside the version space $VS_{H,D}$. Hence, we can restrict our interest to the version space itself.

$$h^* \in \arg \max_{h \in VS_{H,D}} P(h, D)$$

The concept of version space is based on the **inductive learning hypothesis**: any hypothesis that is consistent with the target function over a dataset of *adequate size* will also approximate the target function well over other unobserved examples. In other words, if we consider a dataset of adequate size then every hypotheses inside the version space will be a nice approximation of the target function.

The simplest way to compute the version space is through the **List-Then-Eliminate** algorithm, a brute-force algorithm that enumerates all the hypothesis space and then test the consistency of each hypothesis, discarding the invalid ones. Even though this algorithm is correct, it is also clearly *infeasible* since enumerating all the different hypothesis would require an exponential amount of time. We'll see improved ways to compute the version space given by a dataset.

Algorithm 1: List-Then-Eliminate

Given an hypothesis space H and a dataset D , the algorithm returns $VS_{H,D}$

```

function LISTTHENELIMINATE( $H, D$ )
   $VS_{H,D} := H$ 
  for  $(x, f(x)) \in D$  do
     $H' := \{h \in H \mid h(x) \neq f(x)\}$ 
     $VS_{H,D} = VS_{H,D} - H'$ 
  end for
  return  $VS_{H,D}$ 
end function

```

1.2.1 Representation power and generalization power

In order for the inductive learning hypothesis to hold, the size of the dataset is a **critical** factor: if the hypothesis space is too *powerful* and the search is complete, then the system won't be able to classify new instances, meaning that we have no generalization power.

For instance, consider a generic concept learning problem described by the target function $c : X \rightarrow \{0, 1\}$. Let D be the chosen dataset. For any hypothesis space H , it's easy to see that H is actually *associated* with a particular set of instances, that being all instances that are classified as positive by such hypothesis. In fact, we have a mapping ϕ between the hypothesis space H and the power set $\mathcal{P}(X)$.

$$\phi_H : H \rightarrow \mathcal{P}(X) : h_A \mapsto A = \{x \in X \mid h_A(x) = 1\}$$

In general, this mapping is not surjective, meaning that there is a subset A of $\mathcal{P}(X)$ that is not covered by an hypothesis inside H . In fact, we prefer cases where such

mapping is not surjective. This is known as the **hypothesis space representation issue**: some hypothesis spaces may be useless even when we restrict our interest to the version space.

For example, suppose that there is an hypothesis spaces H_1, H_2 such that H_1 cannot represent $\mathcal{P}(X)$, i.e. ϕ_{H_1} is not surjective, while H_2 can:

- In $VS_{H_1,D}$ we have that $\exists x \in X - X_D$ for which there are two hypotheses $h, h' \in VS_{H_1,D}$ such that $h(x) = 0$ and $h'(x) = 1$.
- In $VS_{H_2,D}$ we have that $\forall x \in X - X_D$ there are two hypotheses $h, h' \in VS_{H_2,D}$ such that $h(x) = 0$ and $h'(x) = 1$.

The small difference in the quantifier has enormous impacts on the usefulness of these two spaces. If we use H_1 then we expect that the approximation found by any algorithm will give the wrong value for some unlabeled inputs $x \in X - X_D$ due to the presence of two functions that can be chosen for x . If we use H_2 , instead, we expect that every unlabeled input will have a wrong value.

These observations imply that the *more information* the hypothesis space encapsulates about the values in X_D , the *harder* it becomes to **generalize** and predict values for samples outside X_D . In other words, a more expressive hypothesis space can **overfit** to the data, making it more difficult to make accurate predictions on unsampled data. We'll return on the problem of overfitting the dataset in following sections.

The process of reducing the *representation power* of the hypothesis space in favor of *generalization power* – as in reducing the hypothesis space from H' to H in the previous example – is called **language bias**. Ideally, we want our hypothesis space to be as good as possible. Clearly, the best possible hypothesis space contains only the optimal approximating function $h^* \in H$. In this case, the previous observations regarding the hypothesis space representation issue are solved: every unlabeled data will have only one value inside the function. The process of selecting one particular hypothesis among the set of possible ones – i.e., choosing $h^* \in H$ – is called **search bias**.

In machine learning, the concept of learning bias is crucial for improving a model's ability to generalize. A good learning bias helps guide the learning algorithm towards patterns in the data that are useful for predicting unseen samples, increasing the system's generalization power. This bias allows the model to make accurate predictions on new data that wasn't part of the training set. Without such a bias, a system would simply **memorize** the dataset, failing to predict values for samples outside the training set, rendering it *ineffective* in real-world applications. Systems lacking generalization capabilities would be of little use, as they wouldn't be able to provide meaningful predictions beyond the data they were trained on.

In real-world applications, datasets often contain **noise**, which refers to irrelevant or erroneous information that can distort the true underlying patterns in the data. Noise can come from a variety of sources, including measurement errors, data entry mistakes, incomplete data or random fluctuations in the system being studied. This noise complicates the *learning process*, as machine learning models may *struggle* to distinguish between true *signal* and *noise*.

A noisy data-point in a dataset D for a function f can be formulated as a pair $(x_i, y_i) \in D$, where $y_i \neq f(x_i)$. This means that there may be *no consistent hypothesis* with noisy data, i.e. $\text{VS}_{H,D} = \emptyset$. In these scenarios, **statistical methods** must be employed to implement robust algorithms, in order to reduce the noise in the data.

1.3 Performance evaluation

After discussing which hypotheses we're interested in, we'll now focus on evaluating the performance of an hypothesis. To give an intuition, we'll focus on classification problems. Let $f : X \rightarrow Y$ be a classification problem. Let \mathcal{D} be a probability distribution over X , meaning that each element in X has an associated probability of being drawn. The sum of all the probabilities in \mathcal{D} must be 1. Let S be a set of n samples drawn from X according to \mathcal{D} and for which we know the value $f(x)$. By definition, S corresponds to a dataset of n random elements. For any possible hypothesis h returned by a learning algorithm obtained through S , we want to know what is the best estimate of the *accuracy* of h over future instances drawn from the same distribution and what is the *probable error* of this accuracy estimate. First, we have to define two error measures.

Definition 6: True error

Let $f : X \rightarrow Y$ be a classification problem. Given an hypothesis h , the **true error** of h with respect to f and a probability distribution \mathcal{D} over X , written as $\text{error}_{\mathcal{D}}(h)$, is defined as:

$$\text{error}_{\mathcal{D}}(h) = \Pr_{x \in \mathcal{D}} [h(x) \neq f(x)]$$

The *true accuracy* of h is defined as $\text{accuracy}_{\mathcal{D}}(h) = 1 - \text{error}_{\mathcal{D}}(h)$.

The true error of an hypothesis is clearly a perfect error measure. However, it's easy to see that computing the true error of h is infeasible since we would have to compute such probability over all the inputs. To fix this, we consider another type of error measure that is less precise but computable.

Definition 7: Sample error

Let $f : X \rightarrow Y$ be a classification problem. Given an hypothesis h , the **sample error** of h with respect to f and a data sample S over \mathcal{D} , written as $\text{error}_S(h)$, is defined as:

$$\text{error}_S(h) = \frac{1}{n} \sum_{x \in S} \delta(x)$$

where $\delta(x) = 1$ if $h(x) \neq f(x)$ and $\delta(x) = 0$ otherwise. The *sample accuracy* is defined as $\text{accuracy}_S(h) = 1 - \text{error}_S(h)$.

The sample error can be efficiently computed over a small data sample. The goal of a learning system is to be accurate in $h(x)$. However, sample accuracy may often fool us into thinking that our hypothesis is good: if $\text{accuracy}_S(h)$ is very high but $\text{accuracy}_{\mathcal{D}}(h)$

is poor, our system is not be very useful.

We also notice that, by definition, $\text{error}_S(h)$ is actually a random variable depending on S . Sampling two different sets S and S' from \mathcal{D} may different values for $\text{error}_S(h)$ and $\text{error}_{S'}(h)$. For this reason, we're also interested in the expected value $\mathbb{E}_{S \subseteq \mathcal{D}}[\text{error}_S(h)]$ of the sample error, i.e. the weighted average over all the possible samples S . The **estimation bias** for \mathcal{D} is defined as the difference between the expected sample error and the true error.

$$\text{bias}_{\mathcal{D}} = \mathbb{E}_{S \subseteq \mathcal{D}}[\text{error}_S(h)] - \text{error}_{\mathcal{D}}$$

We notice that the estimation bias is equal to 0 if and only if $\mathbb{E}_{S \subseteq \mathcal{D}}[\text{error}_S(h)] = \text{error}_{\mathcal{D}}$. However, for any learning algorithm it is *impossible* to prove that the bias is exactly 0. This derives from the very concept of learning function. Hence, in order for an hypothesis to be as good as possible, we want an estimation bias that is as close as possible to 0.

If S is the training set used to compute h through some learning algorithm then $\text{error}_S(h)$ will always be have some bias. In order to get an **unbiased estimate**, the training set S used to compute h and the evaluation set S' must be chosen independently. However, even with an unbiased sample, $\text{error}_{S'}(h)$ may still vary from $\text{error}_{\mathcal{D}}(h)$ – the smaller the set S' , the greater the expected variance.

But how good is an estimate of $\text{error}_{\mathcal{D}}(h)$ provided by a single $\text{error}_S(h)$? From the theory of statistical analysis, we can derive the concept of **confidence interval**. If S contains n samples drawn independently of one another over \mathcal{D} , S is independent of h and $n \geq 30$ then with approximately $N\%$ of probability we have that:

$$|\text{error}_S(h) - \text{error}_{\mathcal{D}}(h)| \leq z_N \sqrt{\frac{\text{error}_S(h)(1 - \text{error}_S(h))}{n}}$$

where the value z_n is a constant given by the *confidence level* $N\%$.

$N\%$	50%	68%	80%	90%	95%	98%	99%
z_n	0.67	1.00	1.28	1.64	1.96	2.33	2.58

Figure 1.1: Relation between each confidence level and its constant

This result shows that a single sample error (when unbiased) is also a good estimation measure for the true error, up to some confidence level. However, in order to get a good approximation we must make some trade offs between *training* and *testing*:

1. Using more samples for training and less for testing gives a better approximating function h , but $\text{error}_S(h)$ may not be a good approximation of $\text{error}_{\mathcal{D}}(h)$.
2. Using less samples for training and more for testing gives a worse approximating function h' , but $\text{error}_S(h')$ is guaranteed to be a good approximation of $\text{error}_{\mathcal{D}}(h)$.

Usually, a good trade off for medium sized datasets is a training set of size $\frac{2}{3} |X|$ and a testing set of size $\frac{1}{3} |X|$. However, computing $\text{error}_S(h)$ is not enough: the estimation bias is defined through $\mathbb{E}_{S \subseteq \mathcal{D}}[\text{error}_S(h)]$. Computing the expected sample error is a task as

hard as computing the true error. Nonetheless, thanks to the **central limit theorem**, the expected sample error can be easily approximated by averaging the computed values of $\text{error}_{S_1}(h), \dots, \text{error}_{S_k}(h)$, where S_1, \dots, S_k are independent from each other.

$$\lim_{k \rightarrow +\infty} \bar{\varepsilon}_k - \mathbb{E}_{S \subseteq \mathcal{D}}[\text{error}_S(h)] = 0 \implies \lim_{k \rightarrow +\infty} \bar{\varepsilon}_k = \mathbb{E}_{S \subseteq \mathcal{D}}[\text{error}_S(h)]$$

where $\bar{\varepsilon}_k = \frac{1}{k} \sum_{i \in [k]} \text{error}_{S_i}(h)$. With all the tools that we have discussed, we're now ready to discuss our first unbiased estimator: the **K-Fold Cross Validation** algorithm.

Algorithm 2: K-Fold Cross Validation for Sample Error

Given a dataset D , a learning algorithm L for a function f and a value $k > 0$, the K-Fold Cross Validation returns an estimation $\text{error}_{L,D}$ of the expected sample error of L over D .

```

function KFOLD_CROSS_VALIDATION( $D, L, k$ )
  Partition  $D$  into  $k$  sets  $S_1, \dots, S_k$  where  $|S_i| > 30$ 
  for  $i = 1, \dots, k$  do
     $T_i = D - S_i$   $\triangleright T_i$  is the training set
     $h_i = L(T_i)$ 
     $\delta_i = \text{error}_{S_i}(h_i)$ 
  end for
  return  $\frac{1}{k} \sum_{i \in [k]} \delta_i$ 
end function

```

1.3.1 Hypothesis comparison

Consider the case where we have two independent hypotheses h_1 and h_2 for some discrete-valued target function. Hypothesis h_1 has been tested on a sample S_1 and h_2 has been tested on a sample S_2 . S_1 and S_2 are independent from each other and they are drawn from the same distribution \mathcal{D} . The difference d between the true errors of these two hypotheses is given by

$$d = \text{error}_{\mathcal{D}}(h_1) - \text{error}_{\mathcal{D}}(h_2)$$

The obvious choice for the estimator \hat{d} of d is given by

$$\hat{d} = \text{error}_{S_1}(h_1) - \text{error}_{S_2}(h_2)$$

In fact, it's very easy to prove that this is indeed an unbiased estimator.

$$\mathbb{E}_{S \subseteq \mathcal{D}}[\text{error}_S(h_1) - \text{error}_S(h_2)] = \mathbb{E}_{S \subseteq \mathcal{D}}[\text{error}_S(h_1)] - \mathbb{E}_{S \subseteq \mathcal{D}}[\text{error}_S(h_2)] = d$$

This definition of error comparison allows us to mathematically describe the concept of **overfitting**. In previous sections, we described this phenomenon as the property of an hypothesis of being too much expressive, i.e. it performs pretty good on the training data but it makes poor predictions on unsampled data.

Definition 8: Overfitting

Let $h \in H$ be an hypothesis for a function f obtained through the training data S , where H is the chosen hypothesis space. We say that h **overfits** the set S if there is an another hypothesis $h' \in H$ obtainable through S that has a sample error higher than h but a true error lower than h .

$$\text{error}_S(h) < \text{error}_S(h') \quad \text{error}_D(h) > \text{error}_D(h')$$

Suppose that we have two learning algorithms L_A and L_B . We want to determine which of them produces the best approximation of f on average. The K-Fold Cross Validation algorithm that we discussed in the previous section can be easily adapted to compare L_A and L_B .

Algorithm 3: K-Fold Cross Validation for Comparison

Given a dataset D , two learning algorithms L_A, L_B for a function f and a value $k > 0$, if the algorithm returns True then we expect that L_A is better than L_B

function KFOLDCROSSCOMPARISON(D, L_A, L_B, k)

Partition D into k sets S_1, \dots, S_k where $|S_i| > 30$

for $i = 1, \dots, k$ **do**

$T_i = D - S_i$

▷ T_i is the training set

$h_i = L_A(T_i)$

$h'_i = L_B(T_i)$

$\delta_i = \text{error}_{S_i}(h_i) - \text{error}_{S_i}(h'_i)$

end for

$\varepsilon = \frac{1}{k} \sum_{i \in [k]} \delta_i$

return $\varepsilon < 0$

end function

1.3.2 Performance metrics

Until now, we have focused on error and accuracy. Intuitively, we expect that if the accuracy is high, i.e. the error is low, then our hypothesis must be a good approximation. However, this is not always the case.

For instance, consider a concept learning problem $f : X \rightarrow \{+, -\}$ with a training set where 90% of the samples are negative, meaning that they are labeled with $-$. Consider the hypothesis h that labels every single element $x \in X_D$ as negative. Even if this hypothesis is clearly bad, the accuracy would still be 90%. Of course, we're usually interested in the accuracy of approximations yield by learning models. In some cases, accuracy only is not enough to assess the performance of a classification method.

In a more statistical sense, error and accuracy can be also described through the concept of positives and negatives:

- A **true positive** is a positive element that gets classified as positive

- A **true negative** is a negative element that gets classified as negative
- A **false positive** is a negative element that gets classified as positive
- A **false negative** is a positive element that gets classified as negative

Here, the error rate is described as the ratio between all the errors (false positives and false negatives) and all the samples:

$$\text{error} = \frac{\text{FP} + \text{FN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

while the accuracy is still defined as the complement of the error.




Figure 1.2: The relevant elements are the elements that are really positive, while the retrieved elements are the ones that have been labeled as positive

As discussed in the example above, these two measures may be inappropriate when the datasets are unbalanced. For this reason, we also consider two additional measures. The first measure is the **recall** (also called *true positive rate*), defined as the ratio between the true positives and the relevant elements. The recall measures the ability of the model to avoid false negatives.

$$\text{recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

The second measure is the **precision** (also called *positive predictive value*), defined as the ratio between the true positives and the retrieved elements. The recall measures the ability of the model to avoid false positives.

$$\text{precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

The recall and the precision are usually used to compute the **F-score**, which measures the predictive performance of a model. More formally, the F_1 -score is the harmonic mean of the precision and recall. Thus, it symmetrically represents both precision and recall in one single metric.

$$F_1 = \frac{2 \cdot \text{recall} \cdot \text{precision}}{\text{recall} + \text{precision}}$$

Another very easy to read measure linked to classification errors is the **confusion matrix**. A confusion matrix reports the percentage of instances of class C_i that have been classified in the class C_j . When the confusion matrix produced by the analysis of the performance of a solution retrieved through an algorithm is “accumulated” towards the central diagonal, the solution is a good model for the problem. Otherwise, the model is considered bad.




Figure 1.3: Example of a good confusion matrix

For regression problems $f : X \rightarrow \mathbb{R}^k$, we need a different type of metrics. In particular, given a test set $S = \{(x_i, t_i)\}_{i \in [n]}$, performance can be measured in various ways:

- **Mean Absolute Error (MAE)**

$$\frac{1}{n} \sum_{i=1}^n \left| \hat{f}(x_i) - t_i \right|$$

- **Mean Squared Error (MSE)**

$$\frac{1}{n} \sum_{i=1}^n (\hat{f}(x_i) - t_i)^2$$

- **Root Mean Squared Error (RMSE)**

$$\sqrt{\frac{1}{n} \sum_{i=1}^n (\hat{f}(x_i) - t_i)^2}$$

The K-Fold Cross Validation algorithm can also be adapted to compute these performance measures.

Algorithm 4: K-Fold Cross Validation for MAE

Given a dataset D , a learning algorithm L for a function f and a value $k > 0$, the K-Fold Cross Validation returns an estimation $\text{MAE}_{L,D}$ of the expected mean absolute error of L over D .

```

function KFOLD_CROSS_VALIDATION( $D, L, k$ )
  Partition  $D$  into  $k$  sets  $S_1, \dots, S_k$  where  $|S_i| > 30$ 
  for  $i = 1, \dots, k$  do
     $T_i = D - S_i$   $\triangleright T_i$  is the training set
     $h_i = L(T_i)$ 
     $\delta_i = \text{MAE}_{S_i}(h_i)$ 
  end for
  return  $\frac{1}{k} \sum_{i \in [k]} \delta_i$ 
end function

```

Classification problems

2.1 Decision trees learning

After discussing the general idea behind machine learning problems and how to evaluate the solution retrieved by an algorithm, we're now ready to focus on how these algorithms work. In particular, we'll start by discussing decision trees.

Given a discrete input space described by m attributes $X = A_1 \times \dots \times A_m$, where each A_i is a finite set, and a problem $f : X \rightarrow Y$, a **decision tree** is a n -ary tree where:

- Each internal node is labeled by an attribute A_i
- Each branch outgoing from a node labeled with A_i denotes a possible value of an attribute $v \in \text{Values}(A_i)$, where the latter is the set of possible values for A_i
- Each leaf node of the tree is labeled with a class $j \in Y$

Given an input, a decision tree computes by querying the internal nodes (starting from the root) and always proceeds on the edge corresponding to the attribute value assumed by the input. In other words, a decision tree is nothing more than a set of rules that classifies any given input. Decision trees are a very easy computational model and they are capable of computing any discrete function. For example, suppose that we want to solve the concept learning problem $\text{PlayTennis} : X \rightarrow \{\text{Yes}, \text{No}\}$ relative to deciding if good conditions are met in order to play tennis. The input set that we'll be working with is the following:

$$X = \{\text{Outlook} \times \text{Temperature} \times \text{Humidity} \times \text{Wind}\}$$

where:

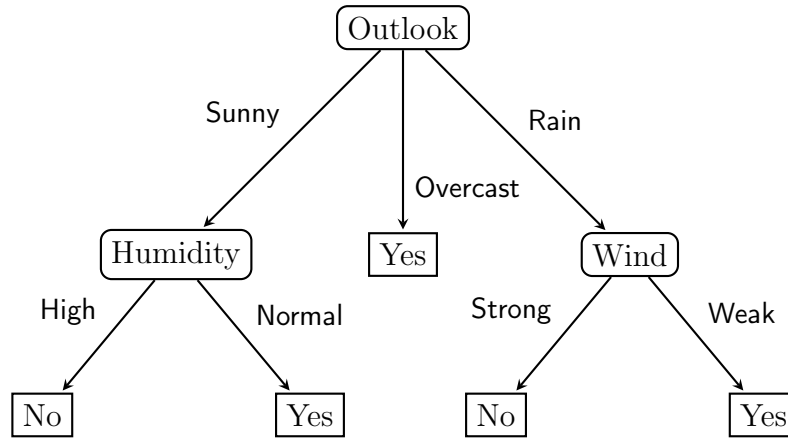
$$\text{Outlook} = \{\text{Sunny}, \text{Overcast}, \text{Rain}\} \quad \text{Temperature} = \{\text{Hot}, \text{Mild}, \text{Cold}\}$$

$$\text{Humidity} = \{\text{Normal}, \text{High}\} \quad \text{Wind} = \{\text{Weak}, \text{Strong}\}$$

Day	Outlook	Temperature	Humidity	Wind	PlayTennis
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

Figure 2.1: Example dataset for *PlayTennis*

A decision tree for this problem would look like the following:

Figure 2.2: Example of decision tree for *PlayTennis*.

In concept learning problems, each decision tree can be described as a disjunction of conjunction of the attributes tested by the tree. In particular, we consider only paths that end up with a positive leaf. For instance, the tree described above can also be represented as following:

$$(\text{Outlook} = \text{Sunny} \wedge \text{Humidity} = \text{Normal})$$

$$\vee (\text{Outlook} = \text{Overcast}) \vee$$

$$(\text{Outlook} = \text{Rain} \wedge \text{Wind} = \text{Weak})$$

2.1.1 Entropy and the ID3 algorithm

In this context, the hypothesis space is the *set of all possible decision trees*. We'll now define the algorithm that we'll be using to derive an optimal decision tree from a given dataset: the **Iterative Dichotomiser 3 (ID3)** algorithm.

Algorithm 5: ID3 algorithm

Given a dataset D for a concept learning problem $f : X \rightarrow \{+, -\}$, an attribute list L containing the attributes of X , the algorithm returns a decision tree for f .

```

function ID3( $D, L$ )
  Initialize an empty decision tree  $T$ 
  Create a root node in  $T$ 
  if All the examples in  $D$  are positive then
    Label the root with +
  else if All the examples in  $D$  are negative then
    Label the root with -
  else if  $L = \emptyset$  then
    Label the root with the most common value for  $A$  in  $D$ 
  else
    Let  $A$  be the best decision attribute in  $L$  for  $D$ 
    Label the root with  $A$ 
    for  $v \in \text{Values}(A)$  do
      Let  $D_a$  be the subset of  $D$  whose tuples have  $A$  set to  $v$ 
      if  $D_a = \emptyset$  then
        Create a leaf node labeled with the most common value for  $A$  in  $D$ 
        Add an edge labeled with  $A = v$  from the root to the leaf node
      else
        Compute the subtree  $T_A = \text{ID3}(D_a, A, L - \{A\})$ 
        Add an edge labeled with  $A = v$  from the root to  $T_A$ 
      end if
    end for
  end if
  Return  $T$ 
end function

```

In the ID3 algorithm described above, the *best decision attribute* is an attribute whose optimality depends on a pre-defined criteria. Based on the criteria chosen, different attributes may be selected. Trying to define an optimal criteria is crucial. In fact, based on the chosen attribute order, we may get a completely different decision tree.

The commonly used measure for this task is the **information gain**, which measures how well a given attribute *separates* the training examples according to their target classification. The information gain of an attribute is measured as reduction in **entropy**, which measures the impurity of a sample.

Definition 9: Entropy

Let S be a sample set for a classification problem $f : X \rightarrow Y$. For each $i \in Y$, we denote with p_i the proportion of elements of S that are classified as C_i . The **entropy** of S is defined as:

$$\text{Entropy}(S) = - \sum_{i \in Y} p_i \log_2 p_i$$

Note: we assume that $0 \log_2 0 = 0$

In the special case of boolean classification problems – since we have only two classes – we denote with p_{\oplus} and p_{\ominus} the proportions of positive and negative samples in S ($p_{\ominus} = 1 - p_{\oplus}$). In this case, the entropy of S is defined as:

$$\text{Entropy} = -p_{\oplus} \log_2 p_{\oplus} - p_{\ominus} \log_2 p_{\ominus}$$




Figure 2.3: The entropy function relative to a boolean classification

Through the graph above, it's easy to see that in boolean classification problems the entropy reaches its maximum value when $p_{\oplus} = 0.5$, i.e. when the two classes are perfectly balanced, while it reaches its minimum when $p_{\oplus} = 0$ or $p_{\oplus} = 1$.

Definition 10: Information gain

The **information gain** for a sample set S is defined as the expected reduction in entropy of S caused by the selection of a value for the attribute A .

$$\text{Gain}(S, A) = \text{Entropy}(S) - \sum_{v \in \text{Values}(A)} \frac{|S_v|}{|S|} \text{Entropy}(S_v)$$

where $S_v = \{s \in S \mid A(s) = v\}$

Suppose that we have a sample set $S = [9+, 5-]$ (where this notation implies that we have 9 positives and 5 negatives) and that we are testing the information gain of the attribute $\text{Wind} = \{\text{Weak}, \text{Strong}\}$. First, we compute the entropy of S :

$$\text{Entropy}(S) = -\frac{9}{14} \log_2 \left(\frac{9}{14} \right) - \frac{5}{14} \log_2 \left(\frac{5}{14} \right) \approx 0.940$$

After looking at the sample set, we get that $S_{\text{Weak}} = [6+, 2-]$ and $S_{\text{Strong}} = [3+, 3-]$. The entropy of these two subsets corresponds to:

$$\text{Entropy}(S_{\text{Weak}}) = -\frac{6}{8} \log_2 \left(\frac{6}{8} \right) - \frac{2}{8} \log_2 \left(\frac{2}{8} \right) \approx 0.811$$

$$\text{Entropy}(S_{\text{Strong}}) = -\frac{3}{6} \log_2 \left(\frac{3}{6} \right) - \frac{3}{6} \log_2 \left(\frac{3}{6} \right) = 1$$

Hence, the information gain on S knowing the attribute Wind corresponds to:

$$\text{Gain}(S, \text{Wind}) = \text{Entropy}(S) - \frac{8}{14} \text{Entropy}(S_{\text{Weak}}) - \frac{6}{14} \text{Entropy}(S_{\text{Strong}}) \approx 0.048$$

Each time the ID3 algorithm has to select the best attribute, it always chooses the one with the highest information gain, corresponding to the one that most reduces the entropy. Eventually, the entropy of the “implicitly partitioned” dataset will reach 1 or 0, meaning that the associated portion of the dataset contains only positive or negative values.

Consider again the dataset S shown in [Figure 2.1](#). This dataset contains 9 positive entries and 5 negative entries. First, the ID3 algorithm computes the information gain of all of the four attributes, finding that Outlook is the attribute that yields the highest gain.

Outlook	Temperature	Humidity	Wind
0.246	0.151	0.048	0.029

Figure 2.4: Gain on S for each attribute

After setting Outlook as the root node, the algorithm now computes the three subtrees, one for each possible value of Outlook . Since the dataset S_{Overcast} contains only positive samples, the leaf node is labeled with Yes. The two other branches, instead, have to recursively compute their subtrees using information gain.




Figure 2.5: The partially learned decision tree after the first level of recursion

After repeating the process for each level of recursion, we get the decision tree shown in Figure 2.2. On each level of recursion, even though such hypothesis space is complete, meaning that every target function lies inside, the ID3 algorithm moves through such space in a greedy manner, returning only a **local minima**, while the target function is a **global minima**. Moreover, each step of the algorithm requires to analyze all the training examples, making this approach **not incremental** – if we want to add more data, the whole tree must be recomputed. However, this also ensures that our statistically-based search choices are robust to noisy data.

2.1.2 Overfitting in decision trees

A common issue in decision tree learning is the **size** of the decision tree yield by the algorithm, that is the number of nodes in the output tree on average. The importance of size comes from the nature of the algorithm itself: since the hypothesis space is complete, if allow the tree to have an huge number of nodes then it will eventually become a perfect approximation of the dataset.




Figure 2.6: Overfitting in decision trees

In the worst case, each entry of the dataset is a leaf of the decision tree. When this happens, the solution found by the algorithm is clearly overfitting the data. In decision tree learning, overfitting is unavoidable. However, it can be reduced through some techniques. The first common technique involves enforcing a maximum growth on the tree by avoiding the recursive process when splitting the data wouldn't give a statistically significant improvement. The second technique grows the full decision tree and then prunes nodes that aren't significant. We'll focus on the second technique.

To determine the optimal tree size, we use a separate set of examples (distinct from the training examples) to evaluate the utility of post-pruning, then apply a statistical test to estimate accuracy of a tree on the entire data distribution.

In **reduced-error pruning**, the data is split into a training and a validation set. The easiest way to achieve each pruning step is to greedily select the subtree that is cut.

1. We copy the decision tree and randomly select a subtree
2. We replace the whole subtree it with a single leaf node labeled with the most common leaf label in the subtree.
3. We test the accuracy on the new subtree through the validation set

After evaluating the accuracy of every possible cut, we choose the one that increases the accuracy the most. Eventually, we'll reach a point where any additional cut will decrease the accuracy, concluding the pruning procedure.




Figure 2.7: Performance after reduced-error pruning.

One more successful method for finding high accuracy hypotheses is the **rule post-pruning** technique. Variants of this technique are used by C4.5, another decision tree learning algorithm.

1. We first infer the decision tree allowing it to overfit the dataset
2. We convert the decision tree into an equivalent set of rules as shown in previous sections

3. We generalize each rule as much as possible independently of others
4. We remove multiple instances of the same generalized rule
5. We sort the final set of rules into a desired sequence
6. We convert the sequence back to a decision tree

Decision trees can also be used for classification of continuous-valued attributes. For instance, given the continuous-valued attribute Temperature, we can create a boolean variable that works with this attribute, such as “Temperature > 72.3”. To work with such variables, the best approach is to pre-define the various splits of the set of continuous values.

Moreover, decision trees can also use multi-valued attributes, such as dates. However, the information gain of these variables is usually too high. To relax this issue, one approach is to use the **Gain Ratio** instead of the gain.

$$\text{GainRatio}(S, A) = \frac{\text{Gain}(S, A)}{\text{SplitInformation}(S, A)}$$

where $\text{SplitInformation}(S, A) = -\sum_{i \in Y} \frac{|S_i|}{|S|} \log_2 \left(\frac{|S_i|}{|S|} \right)$ and $S_v = \{s \in S \mid A(s) = v\}$

In some instances, we may want to give more importance to some particular attributes. To do so, we can define the **cost** for each attribute and replace the gain with one of the following values (they have no standard name)

$$\frac{\text{Gain}^2(S, A)}{\text{Cost}(A)} \quad \frac{2^{\text{Gain}(S, A)} - 1}{(\text{Cost}(A) + 1)^w} \quad \text{for some } w \in [0, 1]$$

If some examples in the dataset have no value for the attribute selected by ID3, we can still use it through one of the following methods:

1. If node u tests A , assign most common value of A among other examples sorted to node u
2. Assign most common value of A among other examples with same target value
3. Assign a probability p_i to each possible value $v_i \in \text{Value}(A)$, assigning fractions of p_i to each descendant in tree

2.2 Bayesian learning

2.2.1 Uncertainty and probability

When the problem that we want to learn is based on actions that may assume continuous values, defining a sequence discrete decision becomes hard. For instance, consider the action A_t representing that Alice leaves for the airport t minutes before her flight. In order to know which t will satisfy the problem, we have to sort out many sub-problems, such as partial observability (road state, other drivers' plans, ...), noisy sensors (traffic reports), uncertainty in action outcomes (flat tire, ...), complexity of modelling and predicting traffic,

Hence an approach purely based on logical deduction will either risk falsehood, leads to conclusions that are too weak for decision making due to them requiring way too many conditions to be met (“ A_{25} get me there on time if there's no accident on the bridge and it doesn't rain and my tires remain intact ...”) or may lead to non-optimal decisions that ensure the outcome (“ A_{1440} will surely suffice, but I have to stay overnight in the airport”). When this is the case, the best option is to just accept **uncertainties** and find a way to work with them. In particular, this implies using a probabilistic approach.

Definition 11: Sample space and probability space

A **sample space** Ω is a set of finite or infinite outcomes. The elements $\omega \in \Omega$ are usually called *atomic event* or *outcome of a random process*.

A **probability space** is a function $\Pr : \Omega \rightarrow [0, 1]$ defined on a sample space where the sum of all the values equals 1

$$\sum_{\omega \in \Omega} \Pr[\omega] = 1$$

Given a sample space Ω , a probability space associates a probability in the range $[0, 1]$ to each element of Ω . Suppose that we want to model a probability space that represents the outcomes of a dice roll. The sample space can be defined as:

$$\Omega = \{1, 2, 3, 4, 5, 6\}$$

If we're using a standard dice, the probability space will be defined as:

$$\Pr[\omega] = \left\langle \frac{1}{6}, \frac{1}{6}, \frac{1}{6}, \frac{1}{6}, \frac{1}{6}, \frac{1}{6} \right\rangle$$

If the dice is loaded, instead, some outcomes may be more-likely. For instance, we could have the following probability space:

$$\Pr[\omega] = \left\langle \frac{1}{2}, 0, \frac{1}{6}, 0, \frac{1}{6}, \frac{1}{6} \right\rangle$$

Definition 12: Event

Given a probability space $\Pr : \Omega \rightarrow [0, 1]$, an **event** is a subset of Ω . The probability of an event $A \subseteq \Omega$ is given by the sum of the probabilities of all of the outcomes in the event.

$$\Pr[A] = \sum_{\omega \in A} \Pr[\omega]$$

Using the previous loaded dice example, the event “roll a value lower than 4” is described by the event $A = \{1, 2, 3\}$ and its probability is $\Pr[A] = \frac{5}{6}$. Events are a simple way to reason about probability. However, sometimes they aren’t as intuitive as they look. For this reason, we usually prefer to work with *random variables*.

Definition 13: Random variable

Given a probability space $\Pr : \Omega \rightarrow [0, 1]$, a **random variable** is a function $X : \Omega \rightarrow B$, where B is an arbitrary set of values.

A random variable associates one of the values in B to every single element of the sample space. To work with a random variable $X : \Omega \rightarrow B$, we often consider the event “ $X = x_i$ ”, where $x_i \in B$. This event is equivalent to the set $\{\omega \in \Omega \mid X(\omega) = x_i\}$. Due to this, a random variable can be viewed both as a function and a variable.

$$\Pr[X = x_i] = \Pr[\{\omega \in \Omega \mid X(\omega) = x_i\}] = \sum_{\substack{\omega \in \Omega : \\ X(\omega) = x_i}} \Pr[\omega]$$

When we’re working with a boolean random variable, i.e. when the value set B is $\{0, 1\}$, we often denote the event “ $X = 1$ ” with X , while the event “ $X = 0$ ” is denoted with $\neg X$. This allows us to write intersections and unions of events as simple conjunctions and disjunctions of boolean random variables.

$$\Pr[\neg A \wedge B] = \sum_{\substack{\omega \in \Omega : \\ A(\omega)=0, B(\omega)=1}} \Pr[\omega]$$

Instead of working with a probability space in order to compute the probabilities of each value assumable by a random variable, we often directly consider a **probability distribution**, a function assigning a probability value to all possible assignments of a random variable. The *joint probability distribution* for a set of random variables gives the probability of every atomic joint event on those random variables.

Pr[PlayTennis \wedge Weather]		Weather			
		Sunny	Rainy	Cloudy	Snowy
PlayTennis	True	0.576	0.02	0.064	0.01
	False	0.144	0.08	0.016	0.09

Figure 2.8: Example of joint probability distribution

We notice that, by definition, the probability of the conjunction of two events is not always equal to the product of the two probabilities of the events.

$$\Pr[A \wedge B] \stackrel{?}{=} \Pr[A] \cdot \Pr[B]$$

When the equality holds, we say that the two events are **independent** from each other. For the probability of the disjunction of two events, instead, we can always use the following formula derived from the *inclusion-exclusion principle*:

$$\Pr[A \vee B] = \Pr[A] + \Pr[B] - \Pr[A \wedge B]$$

Random variables can also influence each others. For instance, if we know that some outcomes of an event A are more likely to happen when some conditions are met, we can restrict our interest to such cases. The probability of an event A given that an event B happened is written $\Pr[A \mid B]$. This is also known as **conditional** (or *posterior*) probability and its defined as:

$$\Pr[A \mid B] = \frac{\Pr[A, B]}{\Pr[B]}$$

where $\Pr[A, B] = \Pr[A \wedge B]$. We notice that the events A and B are independent if and only if $\Pr[A \mid B] = \Pr[A]$. Sometimes, we may be also be given *conditional probability distributions*.

$\Pr[\text{PlayTennis} \mid \text{Weather}]$		Weather			
		Sunny	Rainy	Cloudy	Snowy
PlayTennis	True	0.8	0.2	0.8	0.1
	False	0.2	0.8	0.2	0.9

Figure 2.9: Example of conditional probability distribution

When this is the case, conditional probability can be used to compute joint probability. This is also known as the *product rule*.

$$\Pr[A, B] = \Pr[A \mid B] \cdot \Pr[B]$$

When the values of a random variable Y are mutually exclusive, we may compute the total probability of another random variable X through joint probability or conditional probability. This is called **total probability**.

$$\begin{aligned} \Pr[X = x] &= \Pr[(X = x, Y = y_1) \vee \dots \vee (X = x, Y = y_k)] \\ &= \sum_{i=1}^k \Pr[X = x, Y = y_i] \\ &= \sum_{i=1}^k \Pr[X = x \mid Y = y_i] \cdot \Pr[Y = y_i] \end{aligned}$$

Conditional probability can clearly be viewed as a **normalization factor** α applied to a joint probability. For instance, given the conditional probability, we have that:

$$\Pr[A \mid B] = \frac{\Pr[A, B]}{\Pr[B]} = \alpha \Pr[A, B]$$

where the normalization factor is $\alpha = \frac{1}{\Pr[B]}$. Viewing such probability as nothing more than a normalized instance of the joint one allows us to reason about maximizing and minimizing values in a simpler way: since each element is afflicted by this constant factor, we can just ignore.

$$\arg \max_{x \in \text{Values}(X)} \Pr[X = x \mid Y = y] = \arg \max_{x \in \text{Values}(X)} \alpha \Pr[X = x, Y = y] = \arg \max_{x \in \text{Values}(X)} \Pr[X = x, Y = y]$$

The most important rule derived from the very definition of conditional probability is **Bayes' rule**. This rule allows us to invert the order of the events: to compute the probability of A given B , we can use the probability of B given A

Proposition 1: Bayes' rule

Given two events A and B , it holds that:

$$\Pr[A \mid B] = \frac{\Pr[B \mid A] \cdot \Pr[A]}{\Pr[B]}$$

Bayes' rule becomes even stronger in the context of maximizing and minimizing values thanks to normalization:

$$\arg \max_{x \in \text{Values}(X)} \Pr[X = x \mid Y = y] = \arg \max_{x \in \text{Values}(X)} \Pr[Y = y \mid X = x] \cdot \Pr[X = x]$$

Moreover, when the conditions Y_1, \dots, Y_k are independent from each other, we also get that:

$$\Pr[X \mid Y_1, \dots, Y_k] = \alpha \Pr[Y_1, \dots, Y_k \mid X] \cdot \Pr[X] = \alpha \Pr[Y_1 \mid X] \cdot \dots \cdot \Pr[Y_k \mid X] \cdot \Pr[X]$$

Another interesting consequence of the product rule is the *chain rule*, where we repeatedly apply the product rule on each joint probability:

$$\begin{aligned} \Pr[X_1, \dots, X_{n-1}, X_n] &= \Pr[X_1, \dots, X_{n-1}] \cdot \Pr[X_n \mid X_1, \dots, X_{n-1}] \\ &= \Pr[X_1, \dots, X_{n-2}] \cdot \Pr[X_{n-1} \mid X_1, \dots, X_{n-2}] \cdot \Pr[X_n \mid X_1, \dots, X_{n-1}] \\ &= \prod_{i=1}^n \Pr[X_i \mid X_1, \dots, X_{i-1}] \end{aligned}$$

The chain rule is usually used with **Bayesian networks**, a graphical notation for conditional independence assertions and hence for compact specification of full joint distributions.

Given a set of variables, a Bayesian network is a directed acyclic graph containing one node for each variable. The directed edges of the graph describe influences between variables: an edge (Y, X) implies that Y influences X . Each influence (Y, X) is associated with a conditional probability $\Pr[X \mid Y]$. In the simplest case, the conditional distribution is represented as a **Conditional Probability Table (CPT)** giving the distribution over X for each combination of its parent values.

For instance, consider the following situation. Suppose, that while we're working, our neighbor John calls to say that our alarm is ringing, but our other neighbor Mary doesn't call. However, we know that the alarm can be set off by minor earthquakes. Can we use probabilities to decide if there is a burglar?

First, we define five variables: Burglar, Earthquake, Alarm, JohnCalls, MaryCalls. Then, we consider the relations between such variables based on what we know:

- A burglar can set the alarm, hence $\text{Burglar} \in \text{Parent}(\text{Alarm})$
- An earthquake can set the alarm, hence $\text{Earthquake} \in \text{Parent}(\text{Alarm})$
- The alarm can cause John to call, hence $\text{Alarm} \in \text{Parent}(\text{JohnCalls})$
- The alarm can cause Mary to call, hence $\text{Alarm} \in \text{Parent}(\text{MaryCalls})$


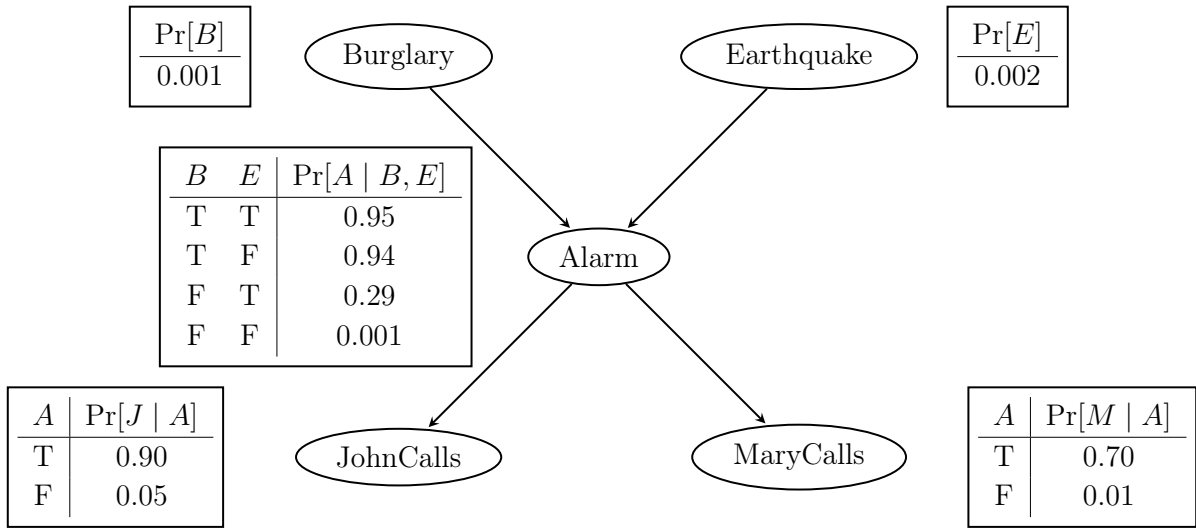


Figure 2.10: The Bayesian network representing the *Burglary Problem*

For each variable, we construct a CPT representing the conditional probabilities for each of the parent attributes.

Figure 2.11: The Bayesian network representing the *Burglary Problem*

Using the chain rule, the probability that both Mary and John call when the alarm rings while there is no burglar or earthquake is given by:

$$\Pr[\neg B, \neg E, A, J, M] = \Pr[\neg B] \cdot \Pr[\neg E] \cdot \Pr[A \mid \neg B, \neg E] \cdot \Pr[J \mid A] \cdot \Pr[M \mid A] \approx 0.00063$$

2.2.2 Optimal Bayes classifier

In classification problems, Bayesian learning provides practical learning algorithms based on multiple parameters, combining prior knowledge (in particular prior probabilities) with new observed data, making probabilistic decisions.

Given a dataset and an instance $x \in X$, the latter gets classified with the most probable class for that instance and that dataset. In other words, given a classification problem $f : X \rightarrow Y$, a dataset D and a new instance $x \in X$, the best class prediction $C^* \in Y$ for x over D is given by:

$$C^* \in \arg \max_{C \in Y} \Pr[C \mid x, D]$$

Similarly, learning is also achieved through probability. Given dataset D and hypothesis space H , we're interested in the hypothesis $h^* \in H$ that maximizes the probability for the given dataset.

$$h^* \in \arg \max_{h \in H} \Pr[h \mid D]$$

As the name suggests, Bayesian learning is based on Bayes' rule. In general, we're interested in two types of hypotheses: the *maximum a posteriori* hypothesis and the *maximum likelihood* hypothesis. The former is achieved by applying Bayes' rule, inverting the order of the events:

$$h_{\text{MAP}} \in \arg \max_{h \in H} \Pr[h \mid D] = \arg \max_{h \in H} \Pr[D \mid h] \cdot \Pr[h]$$

The latter, instead, also assumes that the hypotheses are **uniformly distributed**, i.e. that they all have the same probability. In this case, since $\Pr[h_i] = \Pr[h_j]$ for all $h_i, h_j \in H$, we can ignore this probability:

$$h_{\text{ML}} \in \arg \max_{h \in H} \Pr[h \mid D] = \arg \max_{h \in H} \Pr[D \mid h] \cdot \Pr[h] = \arg \max_{h \in H} \Pr[D \mid h]$$

Definition 14: MAP and ML hypotheses

Given a classification problem $f : X \rightarrow Y$ and a dataset D , the **maximum a posteriori (MAP)** hypothesis h_{MAP} and the **maximum likelihood (ML)** hypothesis h_{ML} are defined as:

$$h_{\text{MAP}} \in \arg \max_{h \in H} \Pr[D \mid h] \cdot \Pr[h]$$

$$h_{\text{ML}} \in \arg \max_{h \in H} \Pr[D \mid h]$$

By definition, h_{MAP}^* and h_{ML}^* are two types of optimal hypotheses for a *given* dataset D . Hence, for any instance $x \in X$, the classes $h_{\text{MAP}}^*(x)$ and $h_{\text{ML}}^*(x)$ may not be the most probable classification for x in *general*, meaning that we don't care about the specific dataset. For instance, suppose we have three hypotheses h_1, h_2, h_3 for a dataset D , where:

$$\Pr[h_1 \mid D] = 0.4 \quad \Pr[h_2 \mid D] = 0.3 \quad \Pr[h_3 \mid D] = 0.3$$

Given an instance $x \in X$, suppose that:

$$h_1(x) = \oplus \quad h_2(x) = \ominus \quad h_3(x) = \ominus$$

Here, for both MAP and ML hypotheses, the most probable hypothesis is clearly h_1 . Hence, we would classify x as type \oplus . However, the most probable class for x in general is clearly \ominus . To fix this issue, we can use the **Optimal Bayes Classifier** method, which is based on total probability. Given an instance $x \in X$ and a class $C \in Y$ with a dataset D , we have that:

$$\Pr[C \mid x, D] = \sum_{h \in H} \Pr[C \mid x, h, D] \cdot \Pr[h \mid x, D] = \sum_{h \in H} \Pr[C \mid x, h] \cdot \Pr[h \mid D]$$

The last equivalence is given by the fact that:

- When h is fixed, the probability that $h(x) = C$ is independent from D , hence $\Pr[C \mid x, h, D] = \Pr[C \mid x, h]$
- h is always independent from x , hence $\Pr[h \mid x, D] = \Pr[h \mid D]$

Definition 15: Optimal Bayes Classifier

Given a classification problem $f : X \rightarrow Y$, a dataset D and a new instance $x \in X - X_D$, the Optimal Bayes Classification for x over D is given by:

$$C_{\text{OB}} \in \arg \max_{C \in Y} \sum_{h \in H} \Pr[C \mid x, h] \cdot \Pr[h \mid D]$$

As the name suggests, this method is an **optimal learner**, meaning that no other classification method using the same hypothesis space and same prior knowledge can outperform this method on average. This method maximizes the probability that the new instance x is classified correctly. For instance, consider again the previous example. We have that:

$$\begin{array}{lll} \Pr[h_1 \mid D] = 0.4 & \Pr[h_2 \mid D] = 0.3 & \Pr[h_3 \mid D] = 0.3 \\ \Pr[\oplus \mid x, h_1] = 1 & \Pr[\oplus \mid x, h_2] = 0 & \Pr[\oplus \mid x, h_3] = 0 \\ \Pr[\ominus \mid x, h_1] = 0 & \Pr[\ominus \mid x, h_2] = 1 & \Pr[\ominus \mid x, h_3] = 1 \end{array}$$

We notice that:

$$\sum_{h_i \in H} \Pr[\oplus \mid x, h_i] \cdot \Pr[h_i \mid D] = 0.4 \qquad \sum_{h_i \in H} \Pr[\ominus \mid x, h_i] \cdot \Pr[h_i \mid D] = 0.6$$

Hence, the Optimal Bayes Classifier would correctly label the instance x with \ominus . To get the full picture behind the power of this optimal learner, we consider a more complex example.

Suppose that we have four kinds of candy bags:

1. The first type contains only candies with the cherry flavor
2. The second type contains 75% of cherry candies and 25% of lime candies
3. The third type contains 50% of cherry candies and 50% of lime candies
4. The fourth type contains 25% of cherry candies and 75% of lime candies
5. The fifth type contains only candies with the lime flavor

A box containing bags of candies arrives to our shop. We know that:

1. 10% of the bags are of the first flavor
2. 20% of the bags are of the second flavor
3. 40% of the bags are of the third flavor
4. 20% of the bags are of the fourth flavor
5. 10% of the bags are of the fifth flavor

We choose a random bag from the box – without knowing its type – and extract some candies from it. We want to learn which type of bag we picked and what is the probability of extracting a candy of a specific flavor next.

First, we model our distribution prior probabilities for the hypothesis space. Each type of bag corresponds to an hypothesis, meaning that:

$$\Pr[H] = \left\langle \frac{1}{10}, \frac{1}{5}, \frac{2}{5}, \frac{1}{5}, \frac{1}{10} \right\rangle$$

The distribution for extracting a lime candy for each hypothesis is given by:

$$\Pr[\ell \mid H] = \left\langle 0, \frac{1}{4}, \frac{1}{2}, \frac{3}{4}, 1 \right\rangle$$

Since we haven't extracted any candy, our previous knowledge is none – hence, the dataset is empty. The probability of the first candy having lime flavor in this case is given by:

$$\Pr[\ell \mid x_1] = \sum_{h_i \in H} \Pr[\ell \mid x, h_i] \cdot \Pr[h_i] = 0 \cdot \frac{1}{10} + \frac{1}{4} \cdot \frac{1}{5} + \frac{1}{2} \cdot \frac{2}{5} + \frac{3}{4} \cdot \frac{1}{5} + 1 \cdot \frac{1}{10} = \frac{1}{2}$$

Suppose that the first candy is indeed a lime candy. For the following extraction, our dataset will be $D_1 = \{(x_1, \ell)\}$. Through Bayes' rule, we have that:

$$\Pr[H \mid D_1] = \frac{\Pr[D_1 \mid H] \cdot \Pr[H]}{\Pr[D_1]} = 2 \Pr[D_1 \mid H] \cdot \Pr[H] = 2 \Pr[\ell \mid H] \cdot \Pr[H]$$

Hence, we have that:

$$\Pr[H \mid D_1] = 2 \cdot \left\langle 0 \cdot \frac{1}{10}, \frac{1}{4} \cdot \frac{1}{5}, \frac{1}{2} \cdot \frac{2}{5}, \frac{3}{4} \cdot \frac{1}{5}, 1 \cdot \frac{1}{10} \right\rangle = \left\langle 0, \frac{1}{10}, \frac{2}{5}, \frac{3}{10}, \frac{1}{5} \right\rangle$$

The probability of extracting a second lime candy is given by:

$$\Pr[\ell \mid x_2, D_1] = \sum_{h_i \in H} \Pr[\ell \mid x_2, h_i] \cdot \Pr[h_i \mid D_1] = 0 \cdot 0 + \frac{1}{4} \cdot \frac{1}{10} + \frac{1}{2} \cdot \frac{2}{5} + \frac{3}{4} \cdot \frac{3}{10} + 1 \cdot \frac{1}{5} = \frac{13}{20}$$

Suppose that we extract another lime candy. The dataset is now $D_2 = \{(x_2, \ell), (x_1, \ell)\}$.

$$\Pr[H \mid D_2] = \frac{\Pr[D_2 \mid H] \cdot \Pr[H]}{\Pr[D_2]} = \frac{20}{13} \Pr[D_2 \mid H] \cdot \Pr[H]$$

Since the data samples in D_2 are independent from each other, we have that:

$$\Pr[H \mid D_2] = \frac{20}{13} \Pr[\{(x_2, \ell)\} \mid H] \cdot \Pr[\{(x_1, \ell)\} \mid H] \cdot \Pr[H] = \frac{20}{13} \Pr[\ell \mid H] \cdot \Pr[H \mid D_1]$$

Again, we now have that:

$$\Pr[H \mid D_2] = \frac{20}{13} \cdot \left\langle 0 \cdot 0, \frac{1}{4} \cdot \frac{1}{10}, \frac{1}{2} \cdot \frac{2}{5}, \frac{3}{4} \cdot \frac{3}{10}, 1 \cdot \frac{1}{5} \right\rangle = \left\langle 0, \frac{1}{26}, \frac{4}{13}, \frac{9}{26}, \frac{4}{13} \right\rangle$$

If we keep extracting a lime candy, the posterior probability of the fifth hypothesis will skyrocket. Using the MAP hypothesis, the fifth one will be selected. However, the selected hypothesis may not be the correct one: the fifth hypothesis clearly overfits our dataset.




Figure 2.12: Probability of the five hypotheses with the increase in the number of samples

Consider now a new manufacturer producing bags with an arbitrary choice of cherry/lime candies. Let $\theta \in [0, 1]$ be the ratio of cherry candies over all the N candies, i.e. $\theta = \frac{\text{num. of cherry candies}}{N}$. For each possible θ , we consider the hypothesis h_θ , where:

$$\Pr[c \mid h_\theta] = \theta \quad \Pr[\ell \mid h_\theta] = 1 - \theta$$

This implies that we have a continuous hypothesis space. The dataset is given by the number of cherry samples c and lime samples ℓ , where $N = c + \ell$. We want to find the ML hypothesis of this setup.

$$h_{\text{ML}}^* \in \arg \max_{\theta \in [0,1]} \Pr[D \mid h_\theta]$$

We notice that:

$$\Pr[D \mid h_\theta] = \prod_{i=1}^N \Pr[d_i \mid h_\theta] = \theta^c \cdot (1 - \theta)^\ell$$

where $d_i \in D$. Hence, we get that:

$$h_{\text{ML}} \in \arg \max_{\theta \in [0,1]} \Pr[D \mid h_\theta] = \theta^c \cdot (1 - \theta)^\ell$$

To get an even better result, we can use the properties of logarithms:

$$h_{\text{ML}} \in \arg \max_{\theta \in [0,1]} \theta^c \cdot (1 - \theta)^\ell = \arg \max_{\theta \in [0,1]} \log(\theta^c \cdot (1 - \theta)^\ell) = \arg \max_{\theta \in [0,1]} c \log \theta + \ell \log(1 - \theta)$$

Since we're working with the continuous interval $[0, 1]$, to find the value $\theta^* \in [0, 1]$ that maximizes the probability we can use derivatives:

$$\frac{dL(D | h_\theta)}{d\theta} = \frac{c}{\theta} - \frac{\ell}{1-\theta} \implies \frac{c}{\theta} - \frac{\ell}{1-\theta} = 0 \implies \theta_{ML} = \frac{c}{c+\ell} = \frac{c}{N}$$

concluding that the value θ_{ML} which gives the best hypothesis h_{ML} is the one representing the ratio of cherry candies in the dataset – as anyone would expect.

In general, given a dataset D where for each $d_i \in D$ we have that $d_i \in \{0, 1\}$, assuming a probability distribution for D over an interval Θ , the maximum likelihood estimation is given by:

$$\theta_{ML} \in \arg \max_{\theta \in \Theta} \log \Pr[d_i | h_\theta]$$

In the particular case of Bernoulli distributions, i.e. where the variable X represents the number of extractions of positive type over N total extractions and for any $k \in [0, N]$ it holds that $\Pr[X = k] = \theta^k (1 - \theta)^{N-k}$, we always get that:

$$\theta_{ML} = \frac{|\{d_i \in D \mid d_i = 1\}|}{|D|}$$

2.2.3 Naïve Bayes classifier

The Bayes Optimal Classifier discussed in the previous section provides the best result, but it's not practical when the hypothesis space is large due to the necessity of enumerating all hypotheses. Instead, we usually work with a weaker version of this classifier, the **Naïve Bayes classifier**.

Consider a target function $f : X \rightarrow Y$. Each instance $x \in X = A_1 \times \dots \times A_n$ is described by a sequence of values a_1, a_2, \dots, a_n , hence:

$$\begin{aligned} \arg \max_{C \in Y} \Pr[C \mid x, D] &= \arg \max_{C \in Y} \Pr[C \mid a_1, \dots, a_n, D] \\ &= \arg \max_{C \in Y} \Pr[a_1, \dots, a_n \mid C, D] \cdot \Pr[C \mid D] \end{aligned}$$

If we assume that the values of a_1, \dots, a_n are independent from each other, we get that:

$$\arg \max_{C \in Y} \Pr[a_1, \dots, a_n \mid C, D] \cdot \Pr[C \mid D] = \arg \max_{C \in Y} \Pr[C \mid D] \prod_{a_i \in x} \Pr[a_i \mid C, D]$$

We observe that the assumption $\Pr[a_1, \dots, a_n \mid C, D] \approx \prod_{a_i \in x} \Pr[a_i \mid C, D]$ is often false. However, it still works surprisingly well, enabling us to estimate the probability without enumerating the whole hypothesis space.

Definition 16: Naïve Bayes Classifier

Given a classification problem $f : X \rightarrow Y$, a dataset D and a new instance $x \in X - X_D$, the Naïve Bayes Classification for x over D is given by:

$$C_{\text{NB}} \in \arg \max_{C \in Y} \Pr[C \mid D] \prod_{a_i \in x} \Pr[a_i \mid C, D]$$

To estimate $\Pr[C \mid D]$, we consider the ratio of positive instances, i.e. instances being classified as class C , over all the instances in the dataset.

$$\widehat{\Pr}[C \mid D] = \frac{|(a, b) \in D \mid b = C|}{|D|}$$

Similarly, to estimate $\Pr[a_i \mid C, D]$ we consider the ratio of positive instances with the value a_i over all the positive instances in the dataset.

$$\widehat{\Pr}[a_i \mid C, D] = \frac{|(a, b) \in D \mid b = C \text{ and } a_i \text{ is an attr. of } a|}{|(a, b) \in D \mid b = C|}$$

Algorithm 6: Naïve Bayes Classifier

Given a dataset D for a classification problem $f : X \rightarrow Y$ and a dataset D , where $X = A_1 \times \dots \times A_k$, the algorithm learns the dataset using the Naïve Bayes Classifier and classifies new instances $x \in X - X_D$.

```

function NAÏVEBAYESLEARN( $X, Y, D$ )
  for  $C \in Y$  do
    Estimate  $\Pr[C \mid D]$  by computing  $\widehat{\Pr}[C \mid D]$ 
    for  $A_k \in X$  do
      for  $a_i \in A_k$  do
        Estimate  $\Pr[a_i \mid C, D]$  by computing  $\widehat{\Pr}[a_i \mid C, D]$ 
      end for
    end for
  end for
end function

function CLASSIFYNEWINSTANCE( $x$ )
   $C_{\text{NB}} \in \arg \max_{C \in Y} \widehat{\Pr}[C \mid D] \prod_{a_i \in x} \widehat{\Pr}[a_i \mid C, D]$ 
end function

```

We notice that this estimate has a flaw: if none of the training instances in the dataset with target class C have attribute value a_i then $\widehat{\Pr}[a_i \mid C, D] = 0$, meaning that $\widehat{\Pr}[C \mid D] \prod_{a_i \in x} \widehat{\Pr}[a_i \mid C, D] = 0$, meaning that no new instance can get classified as C . To fix this issue, we use a small additional value m acting as “virtual” instances and the value p , the previous estimate of $\Pr[a_i \mid C, D]$.

$$\widehat{\Pr}[a_i \mid C, D] = \frac{mp + |(a, b) \in D \mid b = C \text{ and } a_i \text{ is an attr. of } a|}{m + |(a, b) \in D \mid b = C|}$$

We also notice that, since we're interested in the solution maximizing the probability, these approximations don't have to be accurate. In fact, they must only try to preserve which solution will be chose. For instance, the previous estimations preserve the maximum-likelihood solution.

Surprisingly, the Naïve Bayes Classifier can also be adapted to learn how to classify text, i.e. identify them as spam, e-mail, web reviews, etc. Consider the target function $f : \text{Docs} \rightarrow Y$. Any document $d \in \text{Docs}$ is described by a sequence of words $d = w_1 w_2 \dots w_k$. Using the classifier we have that:

$$\Pr[d \mid C, D] = \prod_{w_i \in d} \Pr[w_i \mid C, D]$$

where $\Pr[w_i \mid C, D]$ is the probability of word w_i occurring in a document of class C in D . Let $V = \{w_1, \dots, w_n\}$ be the set of all words appearing in any document $d \in D$. This is referred to as the *vocabulary* of D . After fixing an arbitrary order on V , each document $d \in D$ can be represented by an n -dimensional feature vector over V , where each d_i is an encoding of word w_i in the document, which can be achieved through various methods:

1. **Boolean feature vector:** $d_i = 1$ if w_i occurs in the document, otherwise $d_i = 0$.
2. **Ordinal feature vector:** $d_i = k$ if w_i occurs k times in the document
3. **Real-valued feature vector (tf-idf):** $d_i = \text{tf}(w_i, \text{doc}) \cdot \text{idf}(w_i, D)$, where $\text{tf}(w_i, \text{doc})$ is the frequency of the term w_i in D and $\text{idf}(w_i, \text{doc})$ is the inverse document frequency of w_i in D .

In the first case, we're working with multivariate Bernoulli variables. Hence, we have that:

$$\Pr[d \mid C, D] = \prod_{i=1}^n \Pr[w_i \mid C, D]^{d_i} \cdot (1 - \Pr[w_i \mid C, D])^{1-d_i}$$

For the maximum-likelihood solution, the used estimate of the probability is:

$$\widehat{\Pr}[w_i \mid C, D] = \frac{t_{i,C} + 1}{t_C + 2}$$

where $t_{i,C}$ is the number of documents in D of class C containing w_i and t_C is the number of documents in D of class C . In the second case, instead, we're working with multinomial variables, implying that:

$$\Pr[d \mid C, D] = \frac{n!}{d_1! \cdot \dots \cdot d_n!} \prod_{i=1}^n \Pr[w_i \mid C, D]^{d_i}$$

For the maximum-likelihood solution, the used estimate of the probability is:

$$\widehat{\Pr}[w_i \mid C, D] = \sum_{d \in D} \frac{tf_{i,C} + \alpha}{tf_C + \alpha |V|}$$

where $tf_{i,C}$ is the term frequency of w_i in a document of class C , tf_C is the frequency of all terms in a document of class C and α is a smoothing parameter ($\alpha = 1$ for Laplace smoothing)

This type of text classification has some evident issues. In fact, each of the three methods of representation loses context information – for instance, the order in which the words occur in the document is important. To improve this model, we can eliminate *stop words* (“the”, “a”, etc.), apply *stemming* by replacing words with their basic form (“likes”, “liking” → “like”) or use *n-grams*, where tokens (sequences of words) are used.

2.3 Probabilistic models for classification

2.3.1 Probabilistic generative models

Consider a classification problem $f : \mathbb{R}^d \rightarrow Y$ and a dataset D , we want to estimate $\Pr[C \mid x, D]$ for a new instance $x \in \mathbb{R}^d - \mathbb{R}_D^d$ and a class $C \in Y$. In **probabilistic generative models**, this estimation is achieved through Bayes’ rule, like we did in the previous section for Bayes classifiers. The more general approach for these models is to reduce these probabilities to the computation of a variant of the **sigmoid function**, defined as:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$




Figure 2.13: The sigmoid function

Assuming that $Y = \{C_1, C_2\}$, we can manipulate Bayes’ rule in the following way:

$$\begin{aligned} \Pr[C_1 \mid x, D] &= \frac{\Pr[x \mid C_1, D] \cdot \Pr[C_1 \mid D]}{\Pr[x \mid D]} \\ &= \frac{\Pr[x \mid C_1, D] \cdot \Pr[C_1 \mid D]}{\Pr[x \mid C_1, D] \cdot \Pr[C_1 \mid D] + \Pr[x \mid C_2, D] \cdot \Pr[C_2 \mid D]} \\ &= \frac{1}{1 + \frac{\Pr[x \mid C_2, D] \cdot \Pr[C_2 \mid D]}{\Pr[x \mid C_1, D] \cdot \Pr[C_1 \mid D]}} \end{aligned}$$

By setting $\alpha = \ln \frac{\Pr[x | C_1, D] \cdot \Pr[C_1 | D]}{\Pr[x | C_2, D] \cdot \Pr[C_2 | D]}$, we get the sigmoid function $\sigma(\alpha)$.

Consider now the parametric model where $\Pr[x | C_i, D]$ has the normal distribution $\mathcal{N}(x; \mu_i, \Sigma)$, where μ_i is the mean and Σ is the variance, with the same covariance matrix for both classes. We have that:

$$\alpha = \ln \frac{\Pr[x | C_1, D] \cdot \Pr[C_1 | D]}{\Pr[x | C_2, D] \cdot \Pr[C_2 | D]} = \ln \frac{\mathcal{N}(x; \mu_1, \Sigma) \cdot \Pr[C_1 | D]}{\mathcal{N}(x; \mu_2, \Sigma) \cdot \Pr[C_2 | D]} = w^T x + w_0$$

where:

$$w = \Sigma^{-1}(\mu_1 - \mu_2) \quad w_0 = -\frac{1}{2}\mu_1^T \Sigma^{-1} \mu_1 + \mu_2^T \Sigma^{-1} \mu_2 + \frac{\Pr[C_1 | D]}{\Pr[C_2 | D]}$$

where $\Pr[C_1 | D] = p$ and $\Pr[C_2 | D] = 1 - p$. To summarize, we have that:

$$\Pr[C_1 | x, D] = \sigma(w^T x - w_0) \quad \Pr[C_2 | x, D] = 1 - \sigma(w^T x - w_0)$$

$$\Pr[x | C_1, D] = \mathcal{N}(x; \mu_1, \Sigma) \quad \Pr[x | C_2, D] = \mathcal{N}(x; \mu_2, \Sigma)$$

$$\Pr[C_1 | D] = p \quad \Pr[C_2 | D] = 1 - p$$

This allows us to easily automatize the computation of such probability since we can skip the computation of the logarithm. However, to achieve this we still have to estimate the values of μ_1, μ_2, p and Σ . Given data set $D = \{(x_n, t_n) | i \in [N]\}$, where $t_i = 1$ if x_i is in class C_1 and $t_i = 0$ otherwise. Let N_1 be the number of samples in D belonging to C_1 and N_2 be the number of samples in C_2 ($N = N_1 + N_2$). The maximum-likelihood solution here is given by:

$$\Pr[t | p, \mu_1, \mu_2, \Sigma] = \prod_{i=1}^N (p \mathcal{N}(x; \mu_1, \Sigma))^{t_i} ((1 - p) \mathcal{N}(x; \mu_2, \Sigma))^{1-t_i}$$

After maximizing the logarithmic likelihood, we always obtain that:

$$\begin{aligned} \hat{\mu}_1 &= \frac{1}{N_1} \sum_{i=1}^N t_i x_i & \hat{\mu}_2 &= \frac{1}{N_2} \sum_{i=1}^N (1 - t_i) x_i \\ \hat{p} &= \frac{N_1}{N} & \hat{\Sigma} &= \frac{N_1}{N} S_1 + \frac{N_2}{N} S_2 \end{aligned}$$

where $S_j = \frac{1}{N_j} \sum_{i \in C_j} (x_i - \hat{\mu}_j)(x_i - \hat{\mu}_j)^T$.

More generally, when $Y = \{C_1, \dots, C_k\}$, the dataset is formalized as $D = \{(x_i, t_i) | i \in [N]\}$, where t_i is a one-hot indicator vector with $t_{i,j} = 1$ if and only if x_i is in class C_j . In this case, the estimate is given by:

$$\Pr[C_j | x] = \frac{\Pr[x | C_j, D] \cdot \Pr[C_j, D]}{\sum_{h=1}^j \Pr[x | C_h, D] \cdot \Pr[C_h | D]} = \frac{e^{\alpha_j}}{\sum_{h=1}^j e^{\alpha_h}}$$

where $\alpha_h = \ln \Pr[x | C_h] \Pr[C_h]$. This model is called **Gaussian Naïve Bayes classification**, where $\forall j \in [k]$ we have:

$$\begin{aligned}\widehat{\Pr}[C_j | D] &= \widehat{p}_j = \frac{N_j}{N} & \Pr[x | C_j, D] &= \mathcal{N}(x; \mu_j, \Sigma) \\ \widehat{\mu}_j &= \frac{1}{N_j} \sum_{i=1}^N t_{i,j} x_i & \widehat{\Sigma} &= \sum_{j=1}^k \frac{N_j}{N} S_j\end{aligned}$$

where $S_j = \frac{1}{N_j} \sum_{i \in C_j} (x_i - \widehat{\mu}_j)(x_i - \widehat{\mu}_j)^T$. The use of matrix-like computation can be further extended by also modeling other elements as matrices. For instance, given a problem $f : \mathbb{R}^d \rightarrow Y$, a dataset $D = \{(x_i, t_i) | i \in [N]\}$ can also be represented as $\langle X, t \rangle$, where X is an $N \times d$ matrix of input values and t is a vector of N output values.

To make notation more compact, we set $\tilde{w} = \begin{bmatrix} w_0 \\ w \end{bmatrix}$ and $\tilde{x} = \begin{bmatrix} 1 \\ x \end{bmatrix}$ in order to get:

$$a_h = w^T x + w_0 = \tilde{w}^T \tilde{x}$$

The maximum-likelihood solution for a parametric model M_Θ of dataset $D = \langle X, t \rangle$ with $\Pr[t | \Theta, X]$ is thus given by:

$$\Theta^* \in \arg \max_{\Theta} \ln \Pr[t | \Theta, X]$$

When M_Θ belongs to the exponential family, the probability $\Pr[t | \theta, X]$ can be expressed in the form $\Pr[t | \tilde{w}, X]$

$$\Theta^* \in \arg \max_{\tilde{w}} \ln \Pr[t | \tilde{w}, X]$$

2.3.2 Probabilistic discriminative models

Consider a classification problem $f : \mathbb{R}^d \rightarrow Y$ where $X \subseteq \mathbb{R}^d$ and a dataset $D = \langle X, t \rangle = \{(x_i, t_i) | i \in [N]\}$, with $t_i \in \{0, 1\}$. Differently from generative models, in **probabilistic discriminative models**, to estimate $\Pr[C | x, D]$ for a new instance $x \in X - X_D$ and a class $C \in Y$, we directly estimate the probability:

$$\Pr[C_j | \tilde{x}, D] = \frac{e^{\alpha_j}}{\sum_{h=1}^j e^{\alpha_h}}$$

without using Bayes' rule, where $\alpha_h = \tilde{w}^T \tilde{x}$. In this case, the maximum-likelihood is still given by:

$$\Theta^* \in \arg \max_{\tilde{w}} \ln \Pr[t | \tilde{w}, X]$$

The typical example of discriminative model is **logistic regression**. Suppose that we are working with only two classes. The probability $\Pr[t | \tilde{w}, X]$ – recall that $D = \langle X, t \rangle$ – is given by:

$$\Pr[t | \tilde{w}, X] = \prod_{i=1}^N p_i^{t_i} (1 - p_i)^{1-t_i}$$

where $p_i = \Pr[C_1 \mid \tilde{x}_i] = \sigma(\tilde{w}^T \tilde{x}_i)$. By applying the negative logarithm, we obtain a new type of error function: the **cross-entropy error function**.

Definition 17: Cross-entropy error function

Given a classification problem $f : \mathbb{R}^d \rightarrow Y$ and a dataset $D = \langle X, t \rangle$, the **cross-entropy** (or *negative log likelihood*) error function is defined as:

$$E(\tilde{w}) = -\ln \Pr[t \mid \tilde{w}, X] = -\sum_{i=1}^N t_i \ln p_i + (1 - t_i) \ln(1 - p_i)$$

Logistic regression is based on solving the following optimization problem:

$$\tilde{w}^* \in \arg \min_{\tilde{w}} E(\tilde{w})$$

There are many efficient ways to solve this problem. The most common one is the **iteratively reweighted least squares (IRLS)** method, based on *Newton-Raphson iterative method*, which uses derivatives to get closer and closer to the optimal solution. In particular, since we're working with higher dimensions, we'll be using the concept of gradient and gradient descent.

1. First, compute the gradient of the error with respect to \tilde{w} :

$$\nabla E(\tilde{w}) = \sum_{i=1}^N (p_i - t_i) \tilde{x}_i$$

2. Apply the gradient descent step by setting

$$\tilde{w} := \tilde{w} - H(\tilde{w})^{-1} \nabla E(\tilde{w})$$

where $H(\tilde{w})$ is the Hessian matrix of $E(\tilde{w})$, i.e. $H(\tilde{w}) = \nabla \nabla E(\tilde{w})$.

3. Repeat until a certain improvement threshold is reached or after a fixed number of iterations

Observation 1: Gradient

For those who aren't confident with multivariate calculus, the gradient of a function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is defined as:

$$\nabla f(x) = \begin{bmatrix} \frac{df}{dx_1}(x) \\ \vdots \\ \frac{df}{dx_n}(x) \end{bmatrix}$$

In other words, the gradient of f is the vector whose components are the derivatives of f with respect to the components of the input

These computation can also be made more compact and immediate – thus the strength of this model:

$$\begin{aligned}\nabla E(\tilde{w}) &= \tilde{X}^T(p(\tilde{w}) - t) \\ H(\tilde{w}) &= \tilde{X}^T R(\tilde{w}) \tilde{X}\end{aligned}$$

where $p(\tilde{w}) = [p_1 \ \cdots \ p_N]^T$, $\tilde{X} = [\tilde{x}_1 \ \cdots \ \tilde{x}_N]$ and $R(\tilde{w})$ is the diagonal matrix where $r_{i,i} = p_i(1 - p_i)$.

Algorithm 7: Iterative reweighted least squares (IRLS)

Given a classification problem $f : \mathbb{R}^d \rightarrow Y$, where $Y = C_1, C_2$, and a dataset $D = \langle X, t \rangle$, the algorithm solves the optimization problem $\tilde{w}^* \in \arg \min_{\tilde{w}} E(\tilde{w})$.

```
function IRLS(Y, X, t)
     $\tilde{w} = 0$ 
    do
         $\tilde{w} = \tilde{w} - (\tilde{X}^T R(\tilde{w}) \tilde{X})^{-1} \tilde{X}^T (p(\tilde{w}) - t)$ 
    while  $E(\tilde{w})$  decreases
end function
```

2.4 Linear models for classification

In linear models, we're interested in studying classification problems $f : \mathbb{R}^d \rightarrow Y$ with a dataset whose instances are **linearly separable**. A dataset is said to be linearly separable when there exists an hyperplane that separates the instance space into two regions such that differently classified instances lie on different regions.



Figure 2.14: A linearly separable dataset (left) and a non-linearly separable one (right)

The functions describing such hyperplanes are called **linear discriminant functions**. Suppose that $Y = \{C_1, \dots, C_k\}$. We define k linear functions:

$$y_1(x) = w_1^T x + w_{1_0} = \tilde{w}_1^T \tilde{x} \quad \cdots \quad y_k(x) = w_k^T x + w_{k_0} = \tilde{w}_k^T \tilde{x}$$

where $\tilde{w}_h = \begin{bmatrix} w_{h_0} \\ w_h \end{bmatrix}$ and $\tilde{x} = \begin{bmatrix} 1 \\ x \end{bmatrix}$. Each instance $x \in \mathbb{R}^d$ gets classified as C_j if the $y_j(x)$ is the one with maximum value.

$$C_j \in \arg \max_{C_h \in Y} y_h(x)$$

In particular, we observe that for each $C_j, C_{j'}$ the hyperplane $(\widetilde{w}_j - \widetilde{w}_{j'})^T \widetilde{x} = 0$ acts as a **decision boundary**, splitting the regions of points classified as either C_j or $C_{j'}$.




Figure 2.15: Example of how regions get partitioned by a K -class discriminant

Given the following matrix:

$$y(x) = \begin{bmatrix} y_1(x) \\ \vdots \\ y_k(x) \end{bmatrix} = \begin{bmatrix} \widetilde{w}_1^T \\ \vdots \\ \widetilde{w}_k^T \end{bmatrix} \widetilde{x} = \widetilde{W}^T \widetilde{x}$$

where $\widetilde{W} = [\widetilde{w}_1 \ \cdots \ \widetilde{w}_k]$, learning a linear model is equivalent to estimating the matrix \widetilde{W} for which $y(x) = \widetilde{W}^T \widetilde{x}$ defines the K -class discriminant.

2.4.1 Least squares method

The first method that we'll discuss is the **least squares method**. Consider problem $f : \mathbb{R}^d \rightarrow Y$ a dataset $D = \{(x_i, t_i) \mid i \in [N]\}$, where each t is a 1-of- K encoding, meaning that $t_{i,j} = 1$ if $x_i \in C_j$, otherwise $t_{i,j} = 0$. Let $\widetilde{X} = [\widetilde{x}_1 \ \cdots \ \widetilde{x}_N]^T$ and let $T = [t_1 \ \cdots \ t_N]$. We notice that each row of the product $\widetilde{X}\widetilde{W}$ contains the predicted class for the associated input, while the matrix T contains the real class for the input. When the two matrices are equal, the model is perfect. Hence, we want to minimize the difference between such matrices. To do so, we minimize the **sum-of-squares error function**.

Definition 18: Sum-of-squares error function

Given a linearly separable classification problem $f : \mathbb{R}^d \rightarrow Y$ and a dataset $D = \{(x_i, t_i) \mid i \in [N]\}$, the **sum-of-squares error function** is defined as:

$$E(\widetilde{W}) = \frac{1}{2} \text{tr} \left((\widetilde{X}\widetilde{W} - T)^T (\widetilde{X}\widetilde{W} - T) \right)$$

Note: each entry $a_{i,i}$ on the diagonal of the inner matrix is the double of the square of the difference between x_i 's predicted class and t_i . Hence, the trace of the is equal to the sum-of-squares. The halving factor is just a convention.

This minimization problem can be actually solved in a closed form. In fact, we have that:

$$\widetilde{W} = (\widetilde{X}^T \widetilde{X})^{-1} \widetilde{X}^T T$$

concluding that:

$$y(X) = \widetilde{W}^T \widetilde{X} = T^T (\widetilde{X}^\dagger)^T X$$

where $\widetilde{X}^\dagger = (\widetilde{X}^T \widetilde{X})^{-1} \widetilde{X}^T$ is the *pseudo-inverse* of \widetilde{X} . We observe that this is the first method we have seen where we can actually compute a perfect optimal solution, making this method very easy to use. However, this model still has some problems: the K -class discriminant obtained is not robust to **outliers**, i.e. samples that are very far away from the other samples that are classified with the same class.




Figure 2.16: Solution found for a problem without outliers (left) and solution found for a problem with outliers (right)

2.4.2 Perceptron

Consider a linear classification problem $f : \mathbb{R}^d \rightarrow \{-1, +1\}$. A **perceptron** is a binary classifier that assigns one of two classes to a real-valued input vector. Let $\widetilde{w} = [w_0 \ w_1 \ \dots \ w_d]$ be a weight vector and let $x \in \mathbb{R}^d$. A perceptron is a function $o(x)$ that can be defined in two ways: **thresholded** or **untresholded**. In the untresholded case, we have that $o(x) = \widetilde{w}^T \widetilde{x}$, while in the thresholded case we have that $o(x) = \text{sign}(\widetilde{w}^T \widetilde{x})$, where: The perceptron $o(x)$ given by \widetilde{w} is defined as:

$$\text{sign}(\widetilde{w}^T \widetilde{x}) = \begin{cases} +1 & \text{if } w_0 + w_1 x_1 + \dots + w_d x_d > 0 \\ -1 & \text{otherwise} \end{cases}$$

In both cases, to learn w from a dataset $D = \{(x_i, t_i) \mid i \in [N]\}$, we want to minimize the squared error (*loss function*):

$$E(w) = \frac{1}{2} \sum_{i=1}^N (t_i - o(x))^2$$




Figure 2.17: A thresholded perceptron

Since we're working with real values, the minimization problem can be solved through derivatives and an iterative approach. This is called the **perceptron training rule**.

1. First, compute the derivative of $E(\tilde{w})$ with respect to each component of w_j of \tilde{w}

$$\frac{dE}{dw_j} = \sum_{i=1}^N (t_i - o(x))(-x_{j,i})$$

2. Update each component w_j by setting $\tilde{w}_j := w_j + \Delta w_j$, where:

$$\Delta w_j = -\eta \frac{dE}{dw_j} = \eta \sum_{i=1}^N (t_i - o(x))x_{j,i}$$

with η being a small constant called *learning rate*

3. Repeat until a certain improvement threshold is reached (or after a fixed number of iterations)

The choice of η is critical: if η is small then each iteration will only slightly move the previous linear discriminant, while if η is large each iteration may significantly move the discriminant.

Moreover, when η is small, the final solution will either be very close to the samples of the first class or very close to the samples of the second class. In other words, the solution cannot lie in the middle of the two sample groups.

After learning the optimal \tilde{w} , the predicted class of a new instance $x \in \mathbb{R}^d - \mathbb{R}_D^d$ is directly given by $o(x)$. It can be proven that this procedure will converge if the training data is linearly separable and η is sufficiently small. However, convergence doesn't imply termination, hence the terminating condition. We also notice that the non-fixed iteration termination condition cannot be reached until there are missclassified samples, since otherwise improvement could still be achieved.

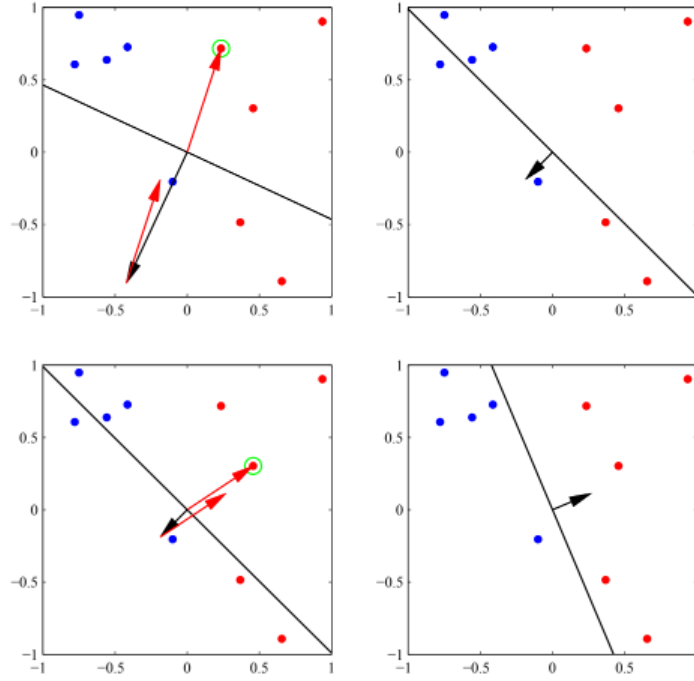


Figure 2.18: Two iterations of the perceptron training rule with a small η value. The vector w is represented by the black arrow, while the vector δw is represented by the red arrow. The green-circled samples are misclassified, hence their feature vector is considered for the next iteration.

2.4.3 Fisher's linear discriminant

Consider a linear classification problem $f : \mathbb{R}^d \rightarrow \{C_1, C_2\}$.