



SAPIENZA
UNIVERSITÀ DI ROMA

“SAPIENZA” UNIVERSITY OF ROME
FACULTY OF INFORMATION ENGINEERING,
INFORMATICS AND STATISTICS
DEPARTMENT OF COMPUTER SCIENCE

Mathematical Logic for Computer Science

Lecture notes integrated with the book “Introduction to
Mathematical Logic”, E. Mendelson

Author
Simone Bianco

March 6, 2025

Contents

Information and Contacts	1
1 Propositional logic	2
1.1 Introduction to logical constructs	2
1.2 Mathematical formalization	4
1.3 Disjunctive and Conjunctive Normal Forms	8
1.4 The compactness theorem	10

Information and Contacts

Personal notes and summaries collected as part of the *Mathematical Logic for Computer Science* course offered by the degree in Computer Science of the University of Rome "La Sapienza".

Further information and notes can be found at the following link:

<https://github.com/Exyss/university-notes>. Anyone can feel free to report inaccuracies, improvements or requests through the Issue system provided by GitHub itself or by contacting the author privately:

- Email: bianco.simone@outlook.it
- LinkedIn: [Simone Bianco](#)

The notes are constantly being updated, so please check if the changes have already been made in the most recent version.

Suggested prerequisites:

Sufficient knowledge of logic and theoretical computer science

Licence:

These documents are distributed under the [GNU Free Documentation License](#), a form of copyleft intended for use on a manual, textbook or other documents. Material licensed under the current version of the license can be used for any purpose, as long as the use meets certain conditions:

- All previous authors of the work must be **attributed**.
- All changes to the work must be **logged**.
- All derivative works must be **licensed under the same license**.
- The full text of the license, unmodified invariant sections as defined by the author if any, and any other added warranty disclaimers (such as a general disclaimer alerting readers that the document may not be accurate for example) and copyright notices from previous versions must be maintained.
- Technical measures such as DRM may not be used to control or obstruct distribution or editing of the document.

1

Propositional logic

1.1 Introduction to logical constructs

Logic is concerned with the validity of reasoning independently of the meaning of its components. In particular, **propositional logic** deals with studying the properties of certain logical constructs used in natural language as well as in scientific and mathematical practice. In particular, we're interested in the following constructs:

- *Negation* (\neg): This operation inverts the truth value of a statement. If a statement is true, its negation is false, and vice versa.
- *Conjunction* (\wedge): This operation represents the logical “and”. The result is true if and only if both components are true.
- *Disjunction* (\vee): This operation represents the logical “or”. The result is true if at least one of the components is true.
- *Implication* (\rightarrow): This operation represents the logical “if ... then” statement. The result is false only if the first component is true and the second component is false. In other words, this operator forces that when the premise holds, the conclusion must always follow (if the premise is false, the conclusion can be either true or false).
- *Equivalence* (\leftrightarrow): This operation represents the logical “if and only if” statement. The result is true when both components have the same truth value, i.e. they're either both true or both false.

For instance, consider the following simple arithmetic argument:

1. If $a = 0$ or $b = 0$, then $a \cdot b = 0$
2. $a \cdot b \neq 0$
3. $a \neq 0$ and $b \neq 0$

Intuitively, the third proposition “follows” from the other two propositions, i.e. it is the conclusion of an argument with the first two as premises. To formalize this argument, we first identify those parts that cannot be further analyzed in terms of logical constructs and that can be true or false (called **atomic parts**). In our case, these atomic parts are $a = 0$, $b = 0$, and $a \cdot b = 0$.

We assign a distinct letter to each of these atomic parts: for $a = 0$, we associate A ; for $b = 0$, we associate B ; and for $a \cdot b = 0$, we associate C . Then, we replace the logical constructs in natural language (if ... then, or, and, not) with formal symbols, called Boolean connectives: \neg for negation, \vee for disjunction, \wedge for conjunction, \rightarrow for implication (if ... then), and \leftrightarrow for equivalence (if and only if).

We obtain the following formalization:

1. $(A \vee B) \rightarrow C$
2. $\neg C$
3. $\neg A \wedge \neg B$

where we read $A \vee B$ as “A or B”, $\neg C$ as “not C” and so on. Consider now the following verbal argument:

1. If the father is tall or the mother is tall, then the child is tall
2. The child is short
3. The father is short and the mother is short

If we attempt a formalization of the argument, we quickly realize that we obtain the same formalization as the previous argument, where we write A for “the father is tall”, B for “the mother is tall”, and C for “the child is tall”. Therefore, the two arguments are identical. Formal logic allows us to identify the common logical structure in arguments that deal with different objects and structures, as is the case in the examples above. However, in this case the first premise is obviously known to be empirically false. Nevertheless, if we assume it as premise for an argument, we intuitively recognize that the reasoning is **valid** (or correct). When we say the argument is valid (or correct), we mean that if the premises are true, then the conclusion is also true; not that the premises are true.

To evaluate the truthfulness of a propositional statement, we use **truth tables**. In logic, a truth table is a table used to determine whether a logical expression is true or false based on the **truth values** of its components. Truth tables are primarily used to evaluate the validity of logical propositions and determine the relationships between various logical statements. Each row of a truth table represents a unique combination of truth values for the variables involved in the logical expression. The columns show the intermediate truth values of the components of the expression, ultimately leading to the evaluation of the entire logical formula.

A	$\neg A$	A	B	$A \wedge B$	A	B	$A \vee B$
0	1	0	0	0	0	0	0
0	1	0	1	0	0	1	1
1	0	1	0	0	1	0	1
		1	1	1	1	1	1

A	B	$A \rightarrow B$	A	B	$A \leftrightarrow B$
0	0	1	0	0	1
0	1	1	0	1	0
1	0	0	1	0	0
1	1	1	1	1	1

Figure 1.1: Truth tables of the five logical operators previously introduces. The value 0 represents the falseness of the propositions, while the value 1 represents their truthfulness.

Consider now a more complex logical expression $A \rightarrow (B \wedge C)$. A truth table for this expression would involve evaluating the truth values for A , B and C , then determining the truth value of the conjunction $B \wedge C$, and finally computing the implication $A \rightarrow (B \wedge C)$.

A	B	C	$B \wedge C$	$A \rightarrow (B \wedge C)$
0	0	0	0	1
0	0	1	0	1
0	1	0	0	1
0	1	1	1	1
1	0	0	0	0
1	0	1	0	0
1	1	0	0	0
1	1	1	1	1

Figure 1.2: Truth table for the logical expression $A \rightarrow (B \wedge C)$

1.2 Mathematical formalization

The first step is to define a completely rigorous formal language that can be subjected to mathematical study. For propositional logic, the procedure is simple: propositions are formalized by formulas obtained by applying Boolean connectives to some atomic building blocks, called *propositional variables*. These variables intuitively represent propositions that cannot be further analyzed in terms of logical operators such as negation, conjunction, disjunction, implication, or equivalence.

Definition 1: Language and propositions

A propositional language is a set $\mathcal{L} = \{p_1, \dots, p_n\}$, where each p_i is a symbol called propositional variable. Given a propositional language \mathcal{L} , the set of propositions (or valid formulas) over \mathcal{L} , denoted with $\text{PROP}_{\mathcal{L}}$ (or $\text{FML}_{\mathcal{L}}$), is the minimal set of finite expressions over $\mathcal{L} \cup \{\neg, \wedge, \vee, \rightarrow, \leftrightarrow\}$ such that:

- Each propositional variable in \mathcal{L} is a proposition, i.e. $\mathcal{L} \subseteq \text{PROP}_{\mathcal{L}}$
- If $A \in \text{PROP}_{\mathcal{L}}$ then $\neg A \in \text{PROP}_{\mathcal{L}}$
- If $A, B \in \text{PROP}_{\mathcal{L}}$ then $(A \wedge B), (A \vee B), (A \rightarrow B), (A \leftrightarrow B) \in \text{PROP}_{\mathcal{L}}$

When \mathcal{L} is non-ambiguous in the context, we denote $\text{PROP}_{\mathcal{L}}$ directly as PROP . Moreover, to make expressions easier to read, we use the following precedence rules: \neg has precedence over \wedge, \vee , while \wedge, \vee have precedence over $\rightarrow, \leftrightarrow$. In other words, we have that $(\neg((\neg A) \vee B)) \rightarrow C$ can be written as $\neg(\neg A \vee B) \rightarrow C$.

After formalizing the concept of propositional language and propositions, we're ready to formalize the concept of truth tables. Each row of a truth table can be described through **assignments**. Here, an assignment is any function $\alpha : \mathcal{L} \rightarrow \{0, 1\}$, i.e. a map that assigns a truth value to each propositional variable of the language. By extension, any assignment over \mathcal{L} induces an assignment over PROP . These assignments are functions $\hat{\alpha} : \text{PROP} \rightarrow \{0, 1\}$ inductively defined as:

$$\hat{\alpha}(A) = \alpha(A) \text{ when } A \in \mathcal{L}$$

$$\hat{\alpha}(\neg A) = \begin{cases} 1 & \text{if } \hat{\alpha}(A) = 0 \\ 0 & \text{if } \hat{\alpha}(A) = 1 \end{cases}$$

$$\hat{\alpha}(A \wedge B) = \begin{cases} 1 & \text{if } \hat{\alpha}(A) = \hat{\alpha}(B) = 1 \\ 0 & \text{otherwise} \end{cases}$$

$$\hat{\alpha}(A \vee B) = \begin{cases} 0 & \text{if } \hat{\alpha}(A) = \hat{\alpha}(B) = 0 \\ 1 & \text{otherwise} \end{cases}$$

$$\hat{\alpha}(A \rightarrow B) = \begin{cases} 0 & \text{if } \hat{\alpha}(A) = 1 \text{ and } \hat{\alpha}(B) = 0 \\ 1 & \text{otherwise} \end{cases}$$

$$\hat{\alpha}(A \leftrightarrow B) = \begin{cases} 1 & \text{if } \hat{\alpha}(A) = \hat{\alpha}(B) \\ 0 & \text{otherwise} \end{cases}$$

To make notation easier, we'll directly refer to $\hat{\alpha}$ as α . Using the concept of assignments, we can also define the concept of **logical equivalence**. Informally, two formulas are said to be logically equivalent when they share the same exact truth table.

Definition 2: Logical equivalence

We say that two formulas $A, B \in \text{PROP}$ are logically equivalent, denoted with $A \equiv B$, when for every assignment α it holds that $\alpha(A) = \alpha(B)$.

By definition, it clearly holds that \equiv is an equivalence relation since it is reflexive, symmetric and transitive. Logical equivalences allow us to treat logical components as if they were algebraic operators. For instance, common algebraic properties include:

- *Involution*: $\neg\neg A \equiv A$
- *Idempotency*: $A \vee A \equiv A$ and $A \wedge A \equiv A$
- *Associativity*: $(A \vee B) \vee C \equiv A \vee (B \vee C)$ and $(A \wedge B) \wedge C \equiv A \wedge (B \wedge C)$
- *Commutativity*: $A \vee B \equiv B \vee A$ and $A \wedge B \equiv B \wedge A$
- *Distributivity*: $A \vee (B \wedge C) \equiv (A \vee B) \wedge (A \vee C)$ and $A \wedge (B \vee C) \equiv (A \wedge B) \vee (A \wedge C)$
- *De Morgan's law*: $\neg(A \vee B) \equiv \neg A \wedge \neg B$ and $\neg(A \wedge B) \equiv \neg A \vee \neg B$

Through these basic properties, we notice that each connective can be rewritten in terms of the other connectives. In fact, we could define propositional logic using only \neg and \wedge (or \neg and \vee). This property is very useful in circuit logic.

- $A \wedge B \equiv \neg(\neg A \vee \neg B)$
- $A \vee B \equiv \neg(\neg A \wedge \neg B)$
- $A \rightarrow B \equiv \neg A \vee B$
- $A \leftrightarrow B \equiv (A \rightarrow B) \wedge (B \rightarrow A)$

If for an assignment α and a formula A it holds that $\alpha(A) = 1$, we say that α *satisfies* A . If A has at least one satisfying assignment, we say that A is **satisfiable**. When no satisfying assignment exists, we say that the formula A is **unsatisfiable**. On the contrary, when every possible assignment satisfies A we say that A is a **tautology** (or *logical truth*, *logically valid*). By definition, tautologies represent statements that are always true, no matter the situation. We denote with SAT, UNSAT and TAUT the sets of satisfiable, unsatisfiable and tautological formulas. By definition, we have that $F \in \text{TAUT}$ if and only if $\neg F \in \text{UNSAT}$. If a formula is not a tautology, we cannot ensure the truthfulness of the statement. Hence, we consider *true statements* only as formulas that are *always true* (the set TAUT), while we consider *false statement* as formulas that are *not always true* (the set SAT – TAUT).

With a slight abuse of notation, we define 1 and 0 as the formulas such that for every assignment it holds that $\alpha(1) = 1$ and $\alpha(0) = 0$. This allows us to further simplify logical equivalences algebraically:

- *Absorption law*: $A \vee 0 \equiv A$ and $A \wedge 1 \equiv A$
- *Cancellation law*: $A \vee 1 \equiv 1$ and $A \wedge 0 \equiv 0$
- *Tertium non datur*: $A \vee \neg A \equiv 1$ and $A \wedge \neg A \equiv 0$

Using all the algebraic rules that we have shown, we can easily evaluate propositional formulas by first reducing them to a smaller equivalent formula and then (eventually) by evaluating the truth table of the reduced form:

$$\begin{aligned}
 A \vee B \rightarrow B \vee \neg C &\equiv \neg(A \vee B) \vee B \vee \neg C \\
 &\equiv (\neg A \wedge \neg B) \vee B \vee \neg C \\
 &\equiv ((\neg A \vee B) \wedge (\neg B \vee B)) \vee \neg C \\
 &\equiv ((\neg A \vee B) \wedge 1) \vee \neg C \\
 &\equiv \neg A \vee B \vee \neg C
 \end{aligned}$$

The concepts of tautology and unsatisfiability are strictly related with the concept of **logical consequence** that we have discussed in the previous section.

Definition 3: Logical consequence

Given the formulas A_1, \dots, A_n, A , we say that A is a logical consequence of A_1, \dots, A_n , written as $A_1, \dots, A_n \models A$ if whenever A_1, \dots, A_n are true A is also true.

Through truth tables, evaluating if we can conclude A through A_1, \dots, A_n is pretty easy: $A_1, \dots, A_n \models A$ if and only if for every row where A_1, \dots, A_n are evaluated as 1 we also have that A is evaluated as 1. Since the number of rows in a truth table is exponential with respect to the number of propositional variables, enumerating the whole truth table is too much expensive. To help with this, the following theorem comes in handy.

Theorem 1

Given the formulas A_1, \dots, A_n, A , the three following statements are equivalent:

1. $A_1, \dots, A_n \models A$
2. $(A_1 \wedge \dots \wedge A_n \rightarrow A) \in \text{TAUT}$
3. $(A_1 \wedge \dots \wedge A_n \wedge \neg A) \in \text{UNSAT}$

Proof. Omitted (trivial) □

Hence, in order to evaluate if $A_1, \dots, A_n \models A$ holds, we can show that $(A_1 \wedge \dots \wedge A_n \rightarrow A) \equiv 1$ or that $(A_1 \wedge \dots \wedge A_n \wedge \neg A) \equiv 0$ through the algebraic properties shown before (this may be as tedious as enumerating the whole truth table, but it is usually faster). Currently, there are no efficient algorithms to answer questions such as “is F satisfiable?”, “is F unsatisfiable?” and “is F a tautology?”. Finding such algorithm or proving that such algorithm cannot exist is equivalent to solving the Millennium Problem $P \stackrel{?}{=} NP$, which asks the question “can every efficiently verifiable problem also be efficiently solved?”. For this problem, the Clay Mathematical Institute offers a prize of one million dollars. In many cases, it is possible to decide whether a certain proposition is a tautology or not, or whether a certain conclusion is a logical consequence of other propositions without constructing the truth table, but only by reasoning rigorously at a higher level.

1.3 Disjunctive and Conjunctive Normal Forms

Every propositional formula can be rewritten in a standard way. In particular, we have two standard ways to write a formula: the **Disjunctive Normal Form (DNF)** and the **Conjunctive Normal Form (CNF)**. A formula is said to be in DNF if it is a disjunction of AND clauses. An AND *clause* is a conjunction of literals. A *literal* is a propositional variable or a negation of a propositional variable. In other words, we say that F is in DNF if:

$$F = \bigvee_{j=1}^m D_j = \bigvee_{j=1}^m (\ell_{1,j} \wedge \ell_{2,j} \wedge \dots \wedge \ell_{k_j,j})$$

where each D_j is an AND clause of the formula and each ℓ_{i_h} is a literal. Similarly, a formula is said to be in CNF if it is a conjunction of OR clauses.

$$F = \bigwedge_{j=1}^m C_j = \bigwedge_{j=1}^m (\ell_{1,j} \vee \ell_{2,j} \vee \dots \vee \ell_{k_j,j})$$

Using the logical equivalence properties, it's clear that every formula has an equivalent DNF and CNF formula. In a more direct way, we can use truth tables (hence assignments) to yield an equivalent DNF or CNF formula.

Theorem 2: Equivalent DNF and CNF formulas

For every formula F there is a DNF formula F^{DNF} and a CNF formula F^{CNF} such that $F \equiv F^{\text{DNF}} \equiv F^{\text{CNF}}$.

Proof. Given $\mathcal{L} = \{p_1, \dots, p_n\}$, consider the truth table of F .

p_1	p_2	\dots	p_n	F
0	\dots	\dots	\dots	\dots
\vdots	\ddots	\ddots	\vdots	\vdots
0	\ddots	\ddots	\vdots	\vdots
1	\ddots	\ddots	\vdots	\vdots
\vdots	\ddots	\ddots	\vdots	\vdots
1	\dots	\dots	\dots	\dots

This truth table has 2^n rows. Each j -th row defines an assignment α_j such that if p_1 assumes value $\alpha_j(p_1)$, p_2 assumes value $\alpha_j(p_2)$ and so on then A assumes value $\alpha_j(A)$. Hence, the cases where A is true are completely described by the rows i where $\alpha_i(A) = 1$.

For each $j \in [2^n]$ and each $i \in [n]$, let $\ell_{i,j}$ be defined as:

$$\ell_{i,j} = \begin{cases} p_i & \text{if } \alpha_j(p_i) = 1 \\ \neg p_i & \text{if } \alpha_j(p_i) = 0 \end{cases}$$

We notice that for every assignment α it holds that:

$$\alpha(\ell_{1,j} \wedge \dots \wedge \ell_{n,j}) = 1 \implies \alpha(\ell_{1,j}) = \alpha_j(p_1), \dots, \alpha(\ell_{n,j}) = \alpha_j(p_n)$$

This implies that for every assignment α it holds that $\alpha(A) = 1$ if and only if for some $j \in [2^n]$ it holds that α corresponds to α_j over p_1, \dots, p_n . Hence, we conclude that $F \equiv F^{\text{DNF}}$, where:

$$F^{\text{DNF}} = \bigvee_{j=1}^{2^n} \bigwedge_{i=1}^n \ell_{i,j}$$

Equivalently, we can obtain F^{CNF} proceeding in the same way by defining each $\ell'_{i,j}$ as:

$$\ell'_{i,j} = \begin{cases} \neg p_1 & \text{if } \alpha_j(p_1) = 1 \\ p_1 & \text{if } \alpha_j(p_1) = 0 \end{cases}$$

In this case, for every assignment α it holds that:

$$\alpha(\ell'_{1,j} \vee \dots \vee \ell'_{n,j}) = 0 \implies \alpha(\ell'_{1,j}) = \alpha_j(p_1), \dots, \alpha(\ell'_{n,j}) = \alpha_j(p_n)$$

This implies that for every assignment α it holds that $\alpha(A) = 0$ if and only if for some $j \in [2^n]$ it holds that α corresponds to α_j over p_1, \dots, p_n . Hence, we conclude that $F \equiv F^{\text{CNF}}$, where:

$$F^{\text{CNF}} = \bigwedge_{j=1}^{2^n} \bigvee_{i=1}^n \ell'_{i,j}$$

□

We observe that the proof of the above theorem is *constructive*, meaning that is based on a procedure that can be used to obtain the DNF and CNF formulas equivalent to an original one. For instance, given the truth table of $A \rightarrow B \wedge C$ shown in [Section 1.1](#), we get the following equivalent DNF formula:

$$(\neg A \wedge \neg B \wedge \neg C) \vee (\neg A \wedge \neg B \wedge C) \vee (\neg A \wedge B \wedge \neg C) \vee (\neg A \wedge B \wedge C) \vee (A \wedge B \wedge C)$$

while the equivalent CNF formula corresponds to:

$$(\neg A \vee B \vee C) \wedge (\neg A \vee B \vee \neg C) \wedge (\neg A \vee B \vee C)$$

We also observe that the second part of the last proof could've also been achieved through a different argument. Consider the fact that $F \equiv \neg \neg F$. Since we know that each formula has a DNF equivalent, we know that there is a formula $(\neg F)^{\text{DNF}}$ such that $\neg F \equiv (\neg F)^{\text{DNF}}$. By iteratively applying De Morgan's law, we can transform $\neg(\neg F)^{\text{DNF}}$ into an equivalent CNF formula. Hence, we conclude that $F^{\text{CNF}} = \neg(\neg F)^{\text{DNF}}$ is the CNF formula such that $F \equiv F^{\text{CNF}}$.

1.4 The compactness theorem

The **compactness theorem** is a fundamental result in mathematical logic first proved by Kurt Gödel. It asserts that if every finite subset of a set of propositions is satisfiable, then the entire set is satisfiable. When the set of propositions is finite, the theorem is trivial. When the set of proposition is countably infinite, i.e. is in bijection with \mathbb{N} , this theorem becomes of particular interest, acting as a **bridge** between finiteness and infiniteness.

Definition 4: Theory

Given a language \mathcal{L} , a theory T over \mathcal{L} is a set of propositions defined on \mathcal{L} , i.e. $T \subseteq \text{PROP}_{\mathcal{L}}$.

A theory is said to be satisfiable when there is an assignment α such that for all $A \in T$ it holds that $\alpha(A) = 1$, otherwise the theory is said to be unsatisfiable – in other words, a theory can be viewed as a conjunction of propositions. With a slight abuse, we extend the notation $\alpha(T)$ to theories for their evaluation.

We observe that while a theory *can* have infinite cardinality every proposition that lies inside it must have a *finite* number of variables since by definition a proposition is a *finite* expression. This fact plays a crucial role for the compactness theorem, which can be stated as follows.

Theorem 3: Compactness theorem

Let T be a theory of finite or countably infinite cardinality and let A be a proposition. Then, A follows from T if and only if there is a finite sub-theory T^{fin} such that A follows from T^{fin} .

$$T \models A \iff \exists T^{fin}_{fin} \subseteq T \text{ such that } T^{fin} \models A$$

As we already mentioned, when the theory T is finite the theorem is trivial. For the countably infinite case, instead, we observe that only one of the two directions is trivial: if A follows from a finite sub-theory $T^{fin}_{fin} \subseteq T$ then A clearly also follows from T since it contains T^{fin} . The other direction appears to be intuitive: in order for A to follow from T , the derivation process must stop “somewhere” in order to prevent an infinite evaluation. All the propositions used in this evaluation process form the finite sub-theory from which A follows. However, such argument is actually a *consequence* of the theorem since we’re guaranteed to stop the evaluation only if the theorem is valid. To prove the other direction, we first define the concept of **finite satisfiability**.

Definition 5: Finite satisfiability

Given a theory T , we say that T is finitely satisfiable when all of its finite subsets are satisfiable.

We observe that the definition of finite satisfiability is different from the definition of satisfiability of a theory. In the former, we require that for every finite sub-theory there is a satisfying assignment, while in the latter we require that there is an assignment that satisfies every finite sub-theory. This inversion of quantifiers is the trick behind the compactness theorem. Before proceeding, we observe that the compactness theorem can actually be restated in an equivalent way.

Proposition 1

Let T be a theory of finite or countably infinite cardinality. Then, the two following points are equivalent:

- $T \models A$ if and only if there is a finite sub-theory T^{fin} such that $T^{fin} \models A$.
- T is satisfiable if and only if it is finitely satisfiable

Proof. Assume that $T \models A$ if and only if $\exists T_0^{fin} \subseteq T$ such that $T_0^{fin} \models A$. Suppose that T is unsatisfiable. Then, this can happen if and only if $T \models 0$. Using the assumption, we get that:

$$T \models 0 \iff \exists T_0^{fin} \subseteq T \text{ } T_0 \models 0$$

Again, $T_0 \models 0$ can happen if and only if T_0 is unsatisfiable, concluding that T is unsatisfiable if and only if $\exists T_0^{fin} \subseteq T$ it holds that T_0 is also unsatisfiable. By contrapositive, we conclude that T is satisfiable if and only if it is finitely satisfiable.

Vice versa, assume now that T is satisfiable if and only if it is finitely satisfiable. If there is a sub-theory $T_0^{fin} \subseteq T$ such that $T_0 \models A$ then $T \models A$ is trivially true. By way of contradiction, suppose that $T \models A$ but $\forall T_0^{fin} \subseteq T$ it holds that $T_0 \not\models A$. Then, we know that $\forall T_0^{fin} \subseteq T$ there is an assignment α such that $\alpha(T_0) = 1$ and $\alpha(A) = 0$, which implies that $\alpha(\neg A) = 1$. Using the assumption, we get that:

$$\forall T_0^{fin} \subseteq T \exists \alpha \text{ } \alpha(T_0) = 1, \alpha(\neg A) = 1 \iff \exists \beta \text{ } \beta(T \cup \{\neg A\}) = 1$$

Since $\beta(T) = 1$ and $\beta(\neg A) = 1$, we conclude that $T \not\models A$. □

Using the above proposition, we can prove the first version of the compactness theorem by proving its second version. Again, one of the two directions of the equivalent statement is trivial: if T is satisfiable then every sub-theory of T must be satisfiable. As in the previous case, the other direction also seems to be intuitive: if T is finitely satisfiable then

there should be a way of combining the assignments $\alpha_1, \alpha_2, \dots$, that satisfy the finite sub-theories T_1, T_2, \dots . However, such combination cannot be easily achieved: two assignments α_i, α_j that satisfy the finite sub-theories T_i, T_j may conflict on a variable, meaning that $\alpha_i(x) \neq \alpha_j(x)$, making their combination not so easy. To prove this direction, we use the following lemma.

Lemma 1: Extendibility of a finitely satisfiable theory

If T is a finitely satisfiable theory then at least one between $T \cup \{A\}$ and $T \cup \{\neg A\}$ is finitely satisfiable.

Proof. By way of contradiction, suppose that T is a finitely satisfiable but both $T \cup \{A\}$ and $T \cup \{\neg A\}$ are not finitely satisfiable. Hence, there are two finite sub-theories $T_0 \subseteq_{fin} T \cup \{\neg A\}$ and $T_1 \subseteq_{fin} T \cup \{A\}$ that are unsatisfiable. Let $\hat{T}_0 = T_0 - \{\neg A\}$ and $\hat{T}_1 = T_1 - \{A\}$. Since $\hat{T}_0 \cup \hat{T}_1 \subseteq T$ and T is finitely satisfiable, there must be an assignment α such that $\alpha(\hat{T}_0 \cup \hat{T}_1) = 1$. Now, we have two cases: if $\alpha(A) = 1$ then $\alpha(T_1) = 1$, while if $\alpha(A) = 0$ then $\alpha(T_0) = 1$. Both cases are a contradiction, concluding the proof. \square

We observe that the lemma is not mutually exclusive. For instance, given the finitely satisfiable theory $T = \{\neg A \vee B, \neg A \vee C\}$, both $T \cup \{\neg A\}$ and $T \cup \{A\}$ are finitely satisfiable. We can now use the lemma to prove the second version of the compactness theorem.

Theorem 4: Compactness theorem (2nd version)

Let T be a theory of finite or countably infinite cardinality and let A be a proposition. Then, T is satisfiable if and only if it is finitely satisfiable.

Proof. The finite case is trivial, hence we focus on the countably infinite case. If T is satisfiable then every sub-theory of T must be satisfiable, including finite ones. Vice versa, suppose that T is finitely satisfiable. Assume that there is an enumeration, i.e. a total order, of $\mathcal{L} = \{p_1, p_2, \dots\}$. Let $T_0 = T$ and for each $i \in \mathbb{N}$ let:

$$T_{i+1} = \begin{cases} T_i \cup \{p_i\} & \text{if } T_i \cup p_i \text{ finitely satisfiable} \\ T_i \cup \{\neg p_i\} & \text{otherwise} \end{cases}$$

By the previous lemma, we know that each T_{i+1} is finitely satisfiable since at least one between $T_i \cup \{p_i\}$ and $T_i \cup \{\neg p_i\}$ must be finitely satisfiable (by construction, when both are finitely satisfiable we give precedence to $T_i \cup \{p_i\}$). Hence, the sequence $T_0 \subseteq T_1 \subseteq \dots$ is well-defined and ensures that they don't conflict with each other.

Claim 1: $T^* = \bigcup_{i \in \mathbb{N}} T_i$ is finitely satisfiable.

Proof of Claim 1. Let X be a finite sub-theory of T^* . By construction, there must be T_i such that $X \subseteq T_i$. Since T_i is finitely satisfiable, X must also be satisfiable. \square

Let α^* be an assignment defined as:

$$\alpha^*(p_i) = \begin{cases} 1 & \text{if } p_i \in T^* \\ 0 & \text{if } \neg p_i \in T^* \end{cases}$$

Claim 2: $\alpha^*(T) = 1$

Proof of Claim 2. Fix $A \in T$. Let p_{i_1}, \dots, p_{i_k} be the variables that appear in A – recall that a proposition always has a finite number of variables. For each $j \in [k]$, let:

$$p_{i_j}^* = \begin{cases} p_{i_j} & \text{if } p_{i_j} \in T^* \\ \neg p_{i_j} & \text{if } \neg p_{i_j} \in T^* \end{cases}$$

Let $A^* = \{A, p_{i_1}^*, \dots, p_{i_k}^*\}$. Since $A^* \subseteq T^*$ and T^* is finitely satisfiable, there must be an assignment β_A such that $\beta_A(A^*) = 1$, which can happen if and only if $\beta_A(A) = 1$ and $\beta_A(p_{i_j}^*) = 1$ for each $j \in [k]$. We notice that by construction it holds that $\beta_A(p_{i_j}^*) = \alpha^*(p_{i_j})$ for any $j \in [k]$. Moreover, we notice that:

- If $p_{i_j} \in T^*$ then $\alpha^*(p_{i_j}) = 1$ and $p_{i_j}^* = p_{i_j}$, implying that $\beta_A(p_{i_j}) = \beta_A(p_{i_j}^*) = 1$ and thus that $\alpha^*(p_{i_j}) = \beta_A(p_{i_j})$
- If $\neg p_{i_j} \in T^*$ then $\alpha^*(p_{i_j}) = 0$ and $p_{i_j}^* = \neg p_{i_j}$, implying that $\beta_A(p_{i_j}) = \neg \beta_A(\neg p_{i_j}) = \neg \beta_A(p_{i_j}^*) = 0$ and thus that $\alpha^*(p_{i_j}) = \beta_A(p_{i_j})$

Since in both cases we have that $\beta(p_{i_1}) = \alpha^*(p_{i_1}), \dots, \beta(p_{i_k}) = \alpha^*(p_{i_k})$ and p_{i_1}, \dots, p_{i_k} are the variables inside A , it must also hold that $\alpha^*(A) = \beta_A(A) = 1$. By applying the same argument on each $A \in T$, we conclude that $\alpha^*(T) = 1$. \square

Since $\alpha^*(T) = 1$, we conclude that T is satisfiable. \square

The Compactness theorem has major implications in mathematics. In particular, it can be applied to prove results in areas outside of logic itself. For instance, it can be used to prove the following three results:

- Every finitely-branching infinite tree has an infinite path (**König's Lemma**)
- A (infinite) graph is k -colorable if and only if every finite subgraph is k -colorable
- Every partial order can be extended to a total order

The idea behind the proofs of the three results is the same. The infinite structure and the property that has to be proven are converted into an infinite theory, which we know to be satisfiable if and only if it is finitely satisfiable, meaning that the statement must be true also for every finite restriction of the infinite case.